

# TRANSACTION

$$\frac{P}{I_1}$$

$\rightarrow I_2 \text{ int } i = S$

| : 

forilux  $I_{60}$

|

|

$I_{n_{100}}$

$$\frac{T_1}{R(A)}$$

$A = A - 10$

$w(A)$

$R(B)$

~~$\times \frac{B = B + 10}{w(B)}$~~

# Transaction

“Transaction is a set of operations which are all logically related.”

OR

“Transaction is a single logical unit of work formed by a set of operations.”

## Operations in Transaction-

The main operations in a transaction are-

- Read Operation
- Write Operation

### 1. Read Operation-

- Read operation reads the data from the database and then stores it in the buffer in main memory.
- For example- **Read(A)** instruction will read the value of A from the database and will store it in the buffer in main memory.

### 2. Write Operation-

- Write operation writes the updated data value back to the database from the buffer.
- For example- **Write(A)** will write the updated value of A from the buffer to the database.

# **ACID Properties-**

- It is important to ensure that the database remains consistent before and after the transaction.
- To ensure the consistency of database, certain properties are followed by all the transactions occurring in the system.
- These properties are called as **ACID Properties** of a transaction.

**A = Atomicity**

**C = Consistency**

**I = Isolation**

**D = Durability**

## 1. Atomicity-

- This property ensures that either the transaction occurs completely or it does not occur at all.
- In other words, it ensures that no transaction occurs partially.
- That is why, it is also referred to as "**All or nothing rule**".
- It is the responsibility of Transaction Control Manager to ensure atomicity of the transactions.

## 2. Consistency-

- This property ensures that integrity constraints are maintained.
- In other words, it ensures that the database remains consistent before and after the transaction.
- It is the responsibility of DBMS and application programmer to ensure consistency of the database.

### **3. Isolation-**

- This property ensures that multiple transactions can occur simultaneously without causing any inconsistency.
- During execution, each transaction feels as if it is getting executed alone in the system.
- A transaction does not realize that there are other transactions as well getting executed parallelly.
- Changes made by a transaction becomes visible to other transactions only after they are written in the memory.
- The resultant state of the system after executing all the transactions is same as the state that would be achieved if the transactions were executed serially one after the other.
- It is the responsibility of concurrency control manager to ensure isolation for all the transactions.

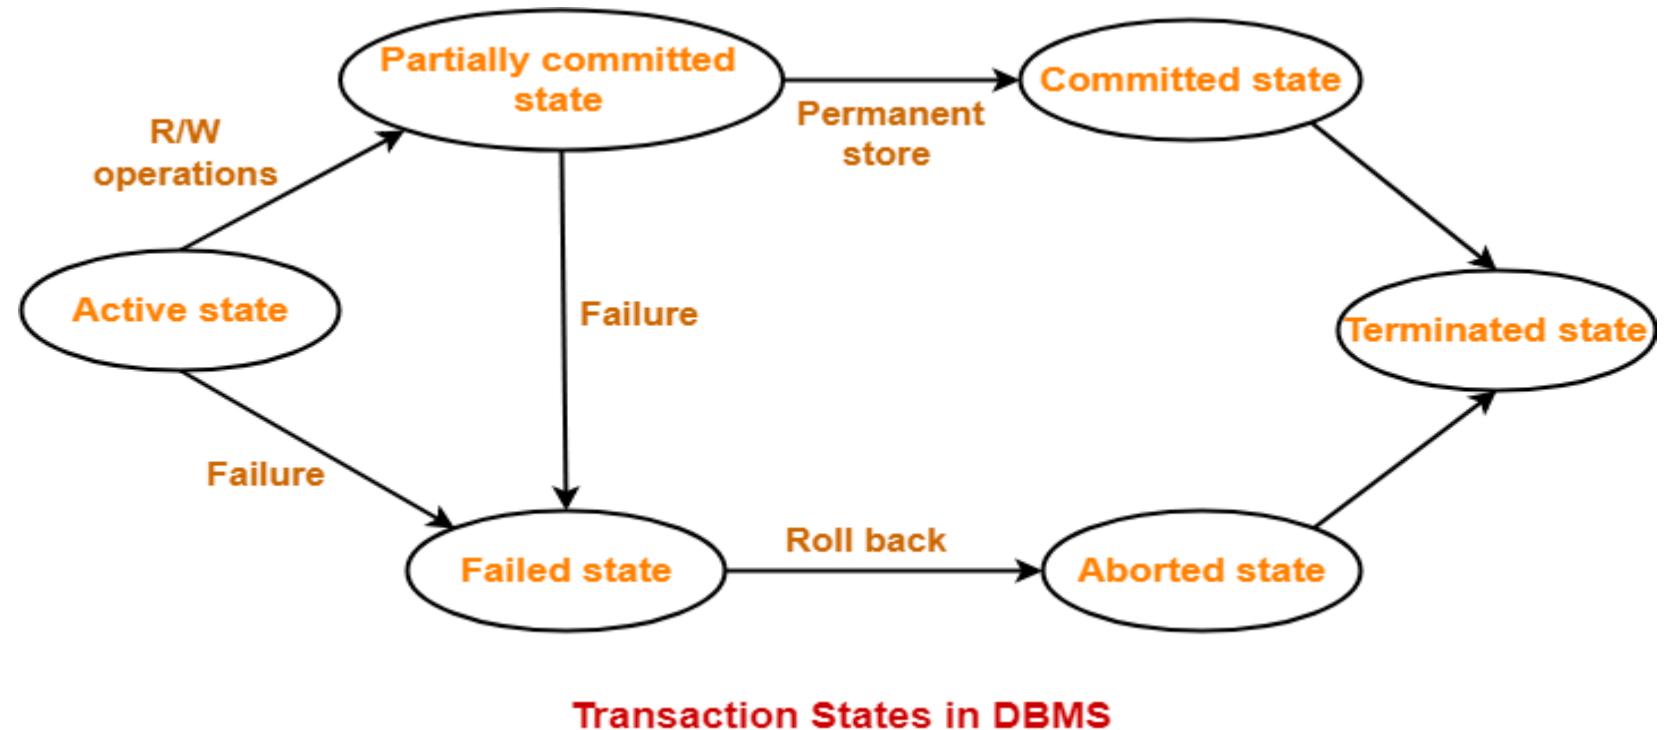
### **4. Durability-**

- This property ensures that all the changes made by a transaction after its successful execution are written successfully to the disk.
- It also ensures that these changes exist permanently and are never lost even if there occurs a failure of any kind.
- It is the responsibility of recovery manager to ensure durability in the database.

# Transaction States-

A transaction goes through many different states throughout its life cycle. These states are called as **transaction states**. Transaction states are as follows-

- Active state
- Partially committed state
- Committed state
- Failed state
- Aborted state
- Terminated state



## 1. Active State-

- This is the first state in the life cycle of a transaction.
- A transaction is called in an **active state** as long as its instructions are getting executed.
- All the changes made by the transaction now are stored in the buffer in main memory.

## 2. Partially Committed State-

- After the last instruction of transaction has executed, it enters into a partially committed state.
- After entering this state, the transaction is considered to be partially committed.
- It is not considered fully committed because all the changes made by the transaction are still stored in the buffer in main memory.

## 3. Committed State-

- After all the changes made by the transaction have been successfully stored into the database, it enters into a **committed state**.
- Now, the transaction is considered to be fully committed.

### NOTE-

- After a transaction has entered the committed state, it is not possible to roll back the transaction.
- In other words, it is not possible to undo the changes that has been made by the transaction.
- This is because the system is updated into a new consistent state.
- The only way to undo the changes is by carrying out another transaction called as **compensating transaction** that performs the reverse operations.

## **4. Failed State-**

- When a transaction is getting executed in the active state or partially committed state and some failure occurs due to which it becomes impossible to continue the execution, it enters into a **failed state**.

## **5. Aborted State-**

- After the transaction has failed and entered into a failed state, all the changes made by it have to be undone.
- To undo the changes made by the transaction, it becomes necessary to roll back the transaction.
- After the transaction has rolled back completely, it enters into an aborted state.

## **6. Terminated State-**

- This is the last state in the life cycle of a transaction.
- After entering the committed state or aborted state, the transaction finally enters into a terminated state where its life cycle finally comes to an end.

# **Concurrency**

Concurrency is the ability of a database to allow multiple users to perform multiple transactions.

## **Advantages of Concurrency**

- **Reduce Waiting Time**
- **Reduce Response Time**
- **Maximize Resource utilization**
- **Increase Efficiency**

# **Concurrency Problems in DBMS-**

- When multiple transactions execute concurrently in an uncontrolled or unrestricted manner, then it might lead to several problems.
- Such problems are called as **concurrency problems**.

The concurrency problems are-

- **Dirty Read Problem**
- **Unrepeatable Read Problem**
- **Lost Update Problem**
- **Phantom Read Problem**

# **1. Dirty Read Problem-**

Reading the data written by an uncommitted transaction is called as dirty read.

This read is called as dirty read because-

- There is always a chance that the uncommitted transaction might roll back later.
- Thus, uncommitted transaction might make other transactions read a value that does not even exist.
- This leads to inconsistency of the database.

## **NOTE-**

- **Dirty read does not lead to inconsistency always.**
- **It becomes problematic only when the uncommitted transaction fails and roll backs later due to some reason.**

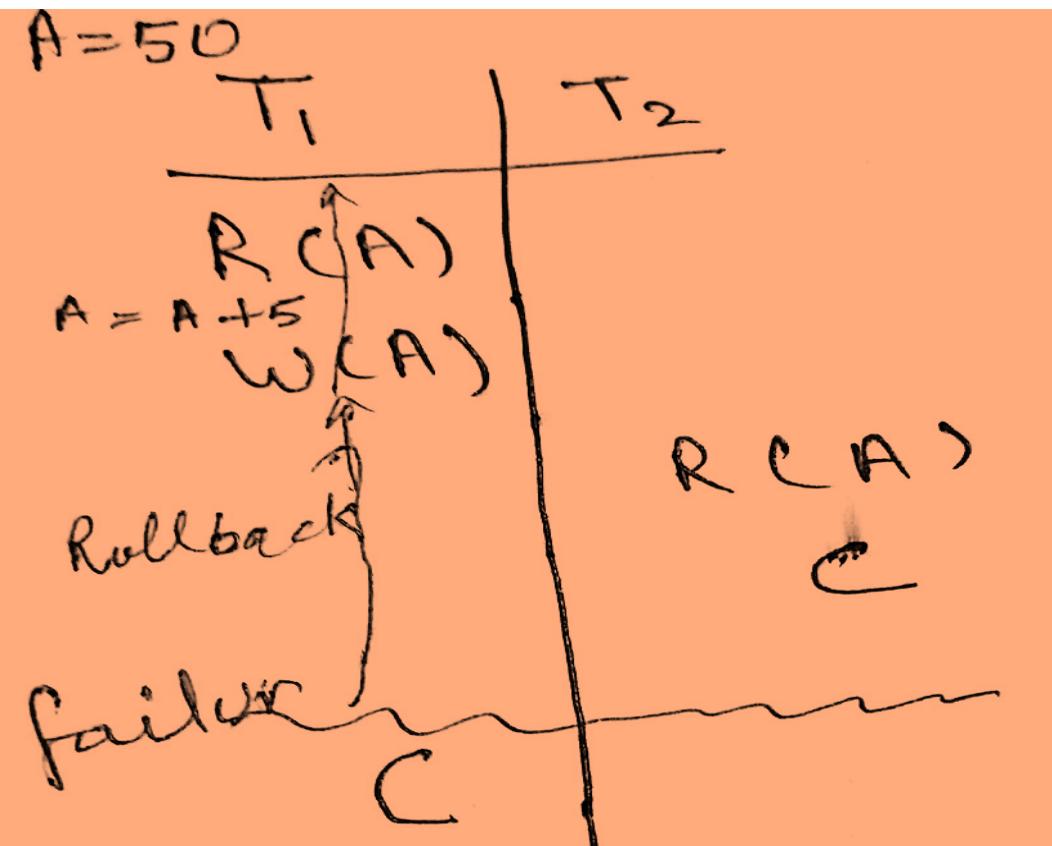
Transaction T1

Transaction T2

R (A)  
W (A)

Failure

R (A) // Dirty Read  
W (A)  
Commit



Here,

T1 reads the value of A.

T1 updates the value of A in the buffer.

T2 reads the value of A from the buffer.

T2 writes the updated the value of A.

T2 commits.

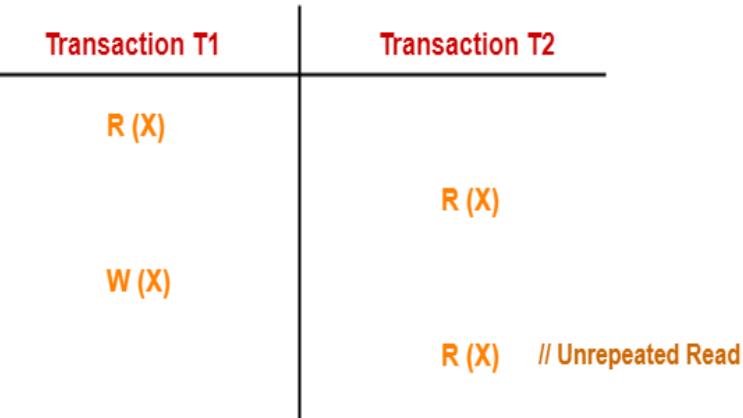
T1 fails in later stages and rolls back.

In this example,

- T2 reads the dirty value of A written by the uncommitted transaction T1.
- T1 fails in later stages and roll backs.
- Thus, the value that T2 read now stands to be incorrect.
- Therefore, database becomes inconsistent.

## 2. Unrepeatable Read Problem-

This problem occurs when a transaction gets to read unrepeated i.e. different values of the same variable in its different read operations even when it has not updated its value.

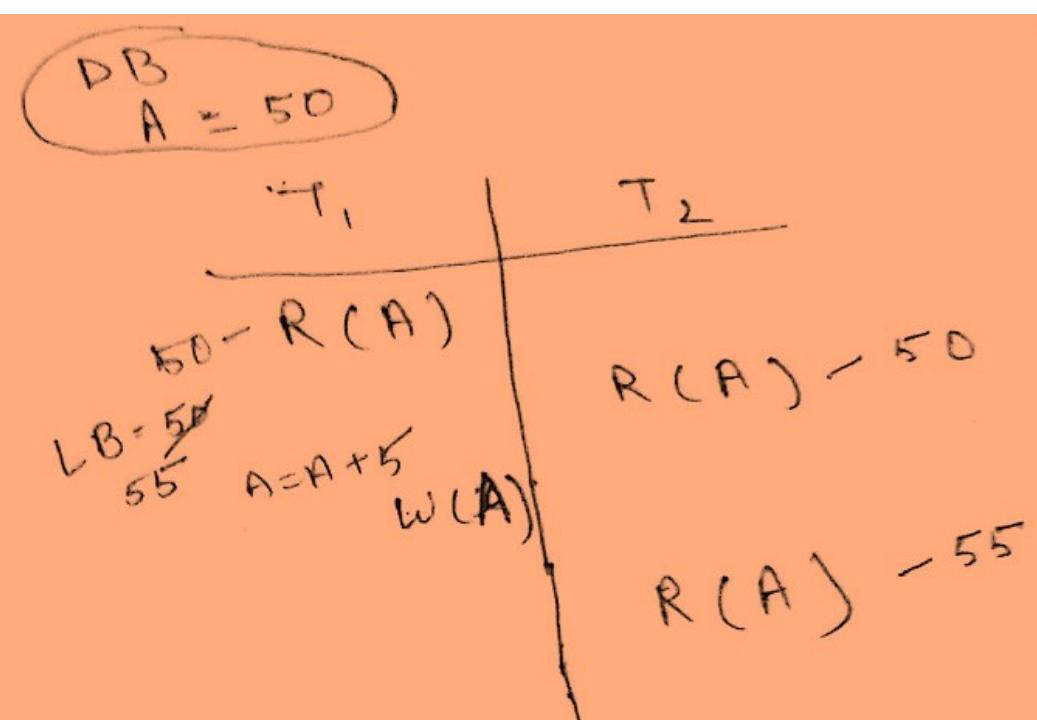


Here,

- T1 reads the value of X (= 10 say).
- T2 reads the value of X (= 10).
- T1 updates the value of X (from 10 to 15 say) in the buffer.
- T2 again reads the value of X (but = 15).

In this example,

- T2 gets to read a different value of X in its second reading.
- T2 wonders how the value of X got changed because according to it, it is running in isolation.



### 3. Phantom Read Problem-

This problem occurs when a transaction reads some variable from the buffer and when it reads the same variable later, it finds that the variable does not exist.

Transaction T1	Transaction T2
R (X)	
Delete (X)	R (X)

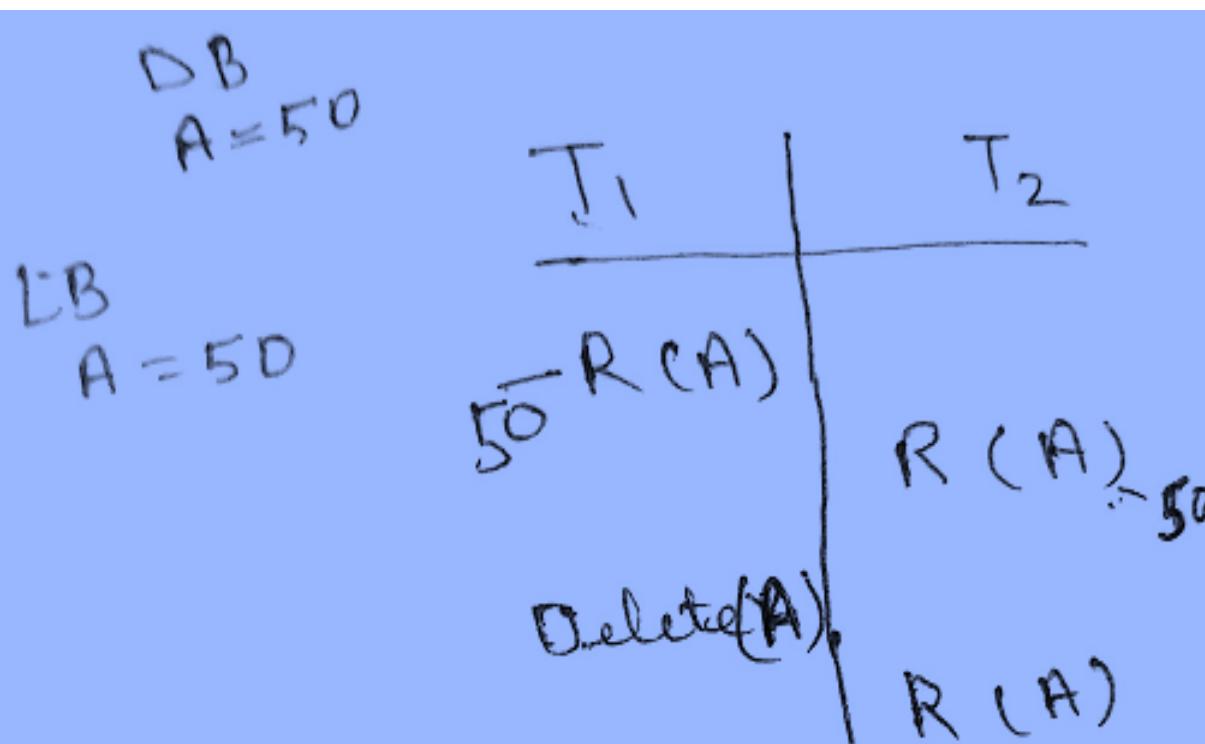
Here,

T1 reads X.

T2 reads X.

T1 deletes X.

T2 tries reading X but does not find it.

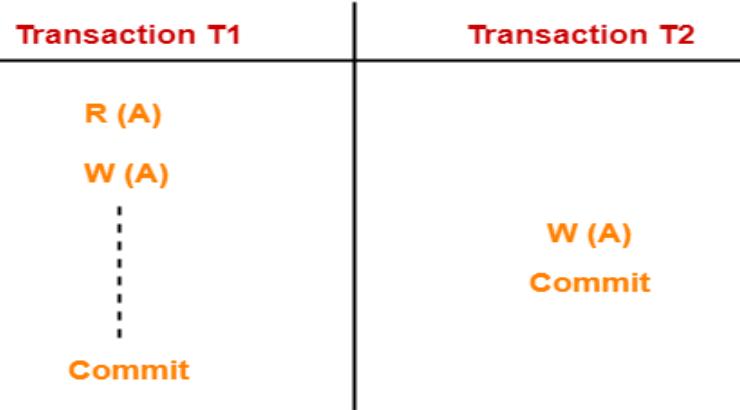


In this example,

- T2 finds that there does not exist any variable X when it tries reading X again.
- T2 wonders who deleted the variable X because according to it, it is running in isolation.

## 4. Lost Update Problem (write - write conflict)-

This problem occurs when multiple transactions execute concurrently and updates from one or more transactions get lost.



Here,

T1 reads the value of A (= 10 say).

T2 updates the value to A (= 15 say) in the buffer.

T2 does blind write A = 25 (write without read) in the buffer.

T2 commits.

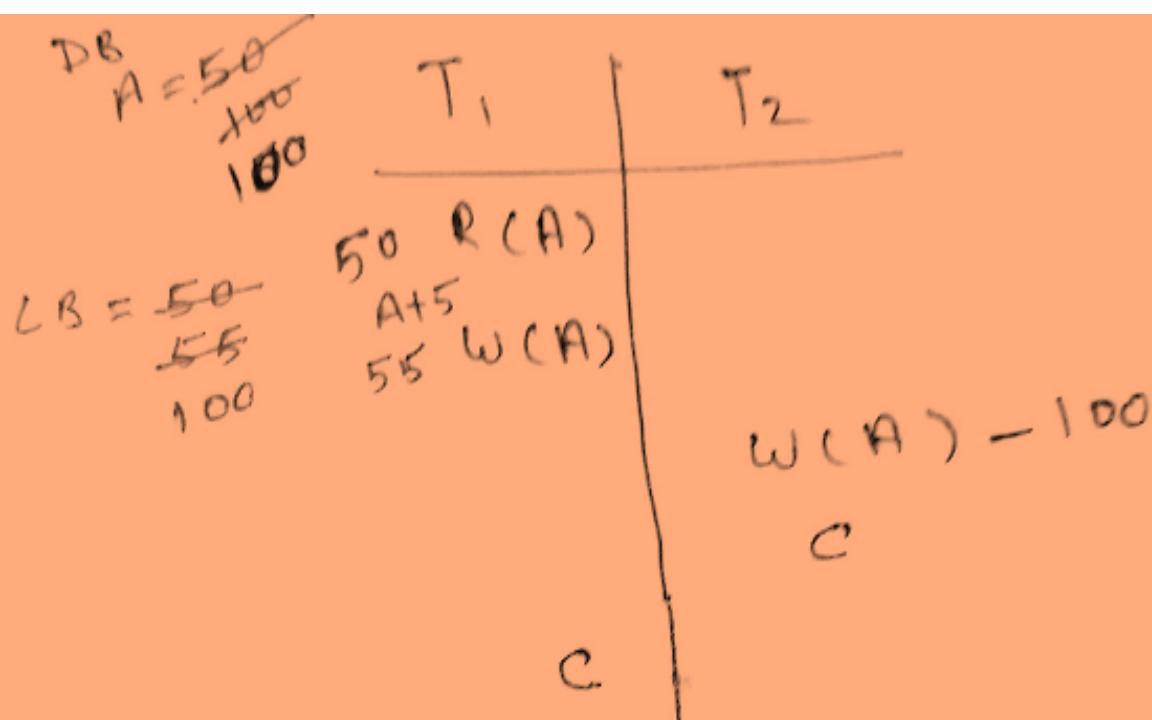
When T1 commits, it writes A = 25 in the database.

In this example,

- T1 writes the over written value of X in the database.
- Thus, update from T1 gets lost.

### NOTE-

- **This problem occurs whenever there is a write-write conflict.**
- **In write-write conflict, there are two writes one by each transaction on the same data item without any read in the middle.**

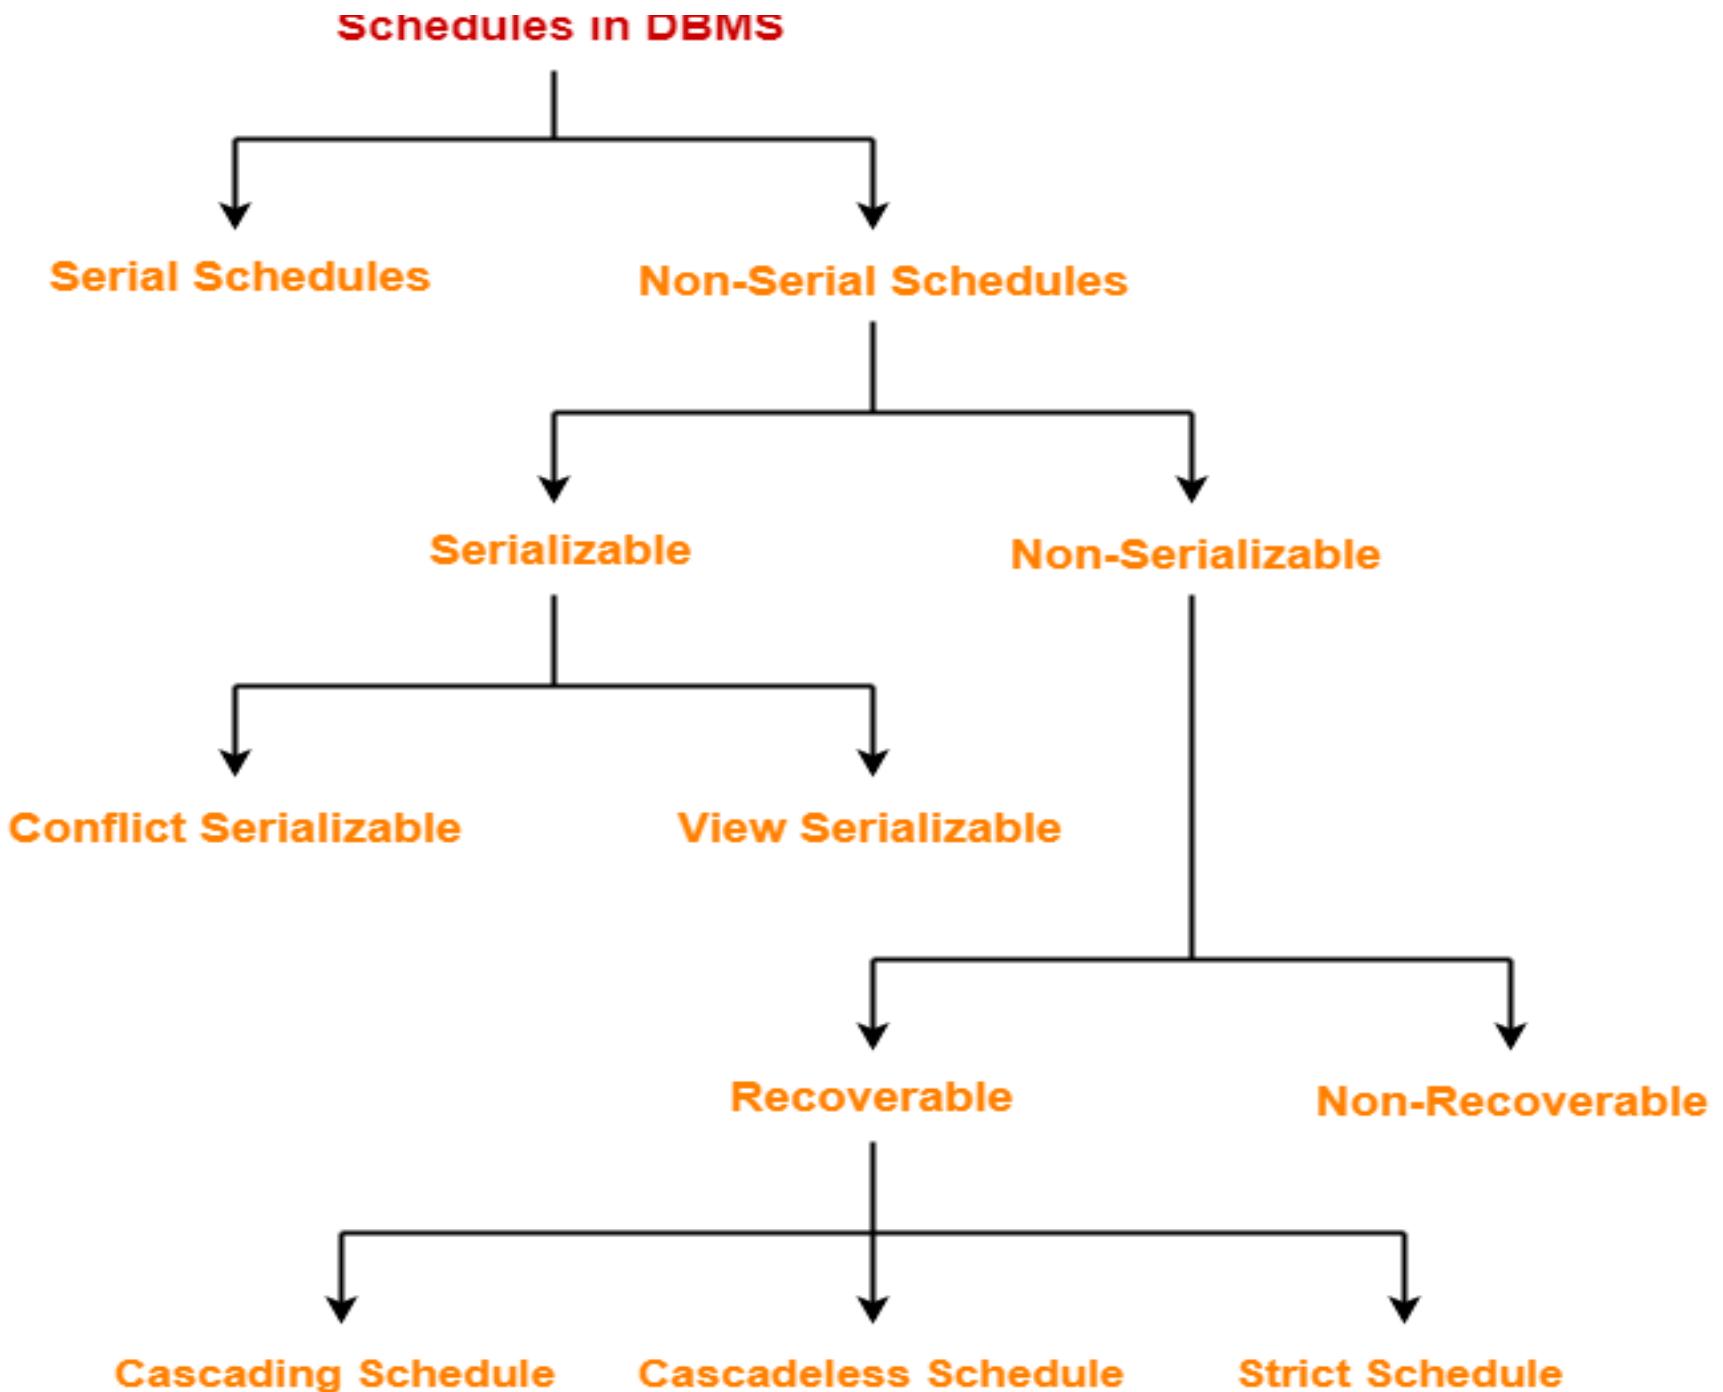


# Avoiding Concurrency Problems-

- To ensure consistency of the database, it is very important to prevent the occurrence of above problems.
- Concurrency Control Protocols help to prevent the occurrence of above problems and maintain the consistency of the database.

# Schedule

The order in which the operations of multiple transactions appear for execution is called as a schedule.



## Serial Schedules-

In serial schedules,

- All the transactions execute serially one after the other.
- When one transaction executes, no other transaction is allowed to execute.

Transaction T1	Transaction T2
R (A)	
W (A)	
R (B)	
W (B)	
Commit	
	R (A)
	W (B)
	Commit

## Characteristics-

Serial schedules are always-

- Consistent
- Recoverable
- Cascadeless
- Strict

Transaction T1	Transaction T2
	R (A)
	W (B)
	Commit
R (A)	
W (A)	
R (B)	
W (B)	
Commit	

## Non-Serial Schedules-

In non-serial schedules,

- Multiple transactions execute concurrently.
- Operations of all the transactions are interleaved or mixed with each other.

	Transaction T1	Transaction T2
	R (A)	
	W (B)	
	R (B)	
	W (B)	
	Commit	
		R (A)
		R (B)
		Commit

## Characteristics-

Non-serial schedules are **NOT** always-

- Consistent
- Recoverable
- Cascadeless
- Strict

	Transaction T1	Transaction T2
	R (A)	
	W (B)	
	R (B)	
	W (B)	
	Commit	
		R (A)
		R (B)
		Commit

$T_1$	$T_2$
$R(A)$	$R(B)$
$w(A)$	$w(B)$
$R(B)$	$R(A)$
$w(B)$	$w(A)$

Schedule :-

$S_1$	
$T_1$	$T_2$
	$R(B)$
	$w(B)$
	$R(A)$
	$w(A)$
$R(A)$	
$w(A)$	
$R(B)$	
$w(B)$	

$S_2$	
$T_1$	$T_2$
$R(A)$	
$w(A)$	
$R(B)$	$R(B)$
$w(B)$	$R(A)$
	$w(A)$

## Finding Number Of Schedules-

Consider there are n number of transactions T<sub>1</sub>, T<sub>2</sub>, T<sub>3</sub> ..., T<sub>n</sub> with N<sub>1</sub>, N<sub>2</sub>, N<sub>3</sub> ..., N<sub>n</sub> number of operations respectively.

## Total Number of Schedules-

Total number of possible schedules (serial + non-serial) is given by-

$$N_1 + N_2 + N_3 + \dots + N_n$$

$$N_1! \times N_2! \times N_3! \times \dots \times N_n!$$

## Total Number of Serial Schedules-

Total number of serial schedules

$$\begin{aligned} &= \text{Number of different ways of arranging } n \text{ transactions} \\ &= n! \end{aligned}$$

## Total Number of Non-Serial Schedules-

Total number of non-serial schedules

$$= \text{Total number of schedules} - \text{Total number of serial schedules}$$

Q. Consider there are three transactions with 2, 3, 4 operations respectively, find-

- How many total number of schedules are possible?
- How many total number of serial schedules are possible?
- How many total number of non-serial schedules are possible?

### Total Number of Schedules-

Using the above formula, we have-

$$\begin{aligned}\text{Total number of schedules} &= \frac{(2 + 3 + 4)!}{2! \times 3! \times 4!} \\ &= 1260\end{aligned}$$

### Total Number of Serial Schedules-

Total number of serial schedules  
= Number of different ways of  
arranging 3 transactions  
=  $3!$   
= 6

### Total Number of Non-Serial Schedules-

Total number of non-serial schedules  
= Total number of schedules – Total number of serial  
schedules  
=  $1260 - 6$   
= 1254

## Serializability in DBMS-

- Some non-serial schedules may lead to inconsistency of the database.
- Serializability is a concept that helps to identify which non-serial schedules are correct and will maintain the consistency of the database.

## Serializable Schedules-

If a given non-serial schedule of 'n' transactions is equivalent to some serial schedule of 'n' transactions, then it is called as a **Serializable schedule**.

## Characteristics-

Serializable schedules behave exactly same as serial schedules.

Thus, serializable schedules are always-

- Consistent
- Recoverable
- Cascadeless
- Strict

## Types of Serializability-

Serializability is mainly of two types-

- Conflict Serializability
- View Serializability

## Conflict Serializability-

If a given non-serial schedule can be converted into a serial schedule by swapping its non-conflicting operations, then it is called as a **conflict serializable schedule**.

## Conflicting Operations-

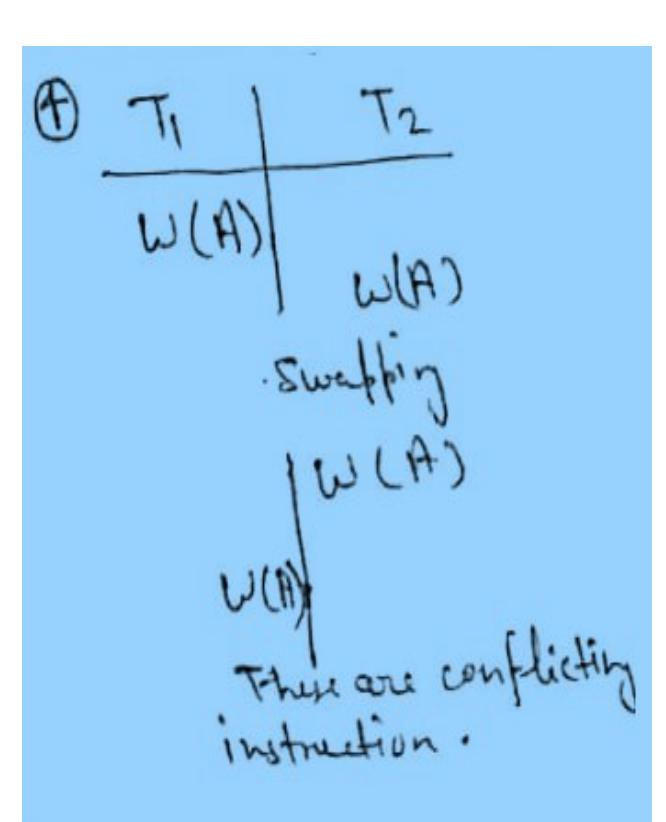
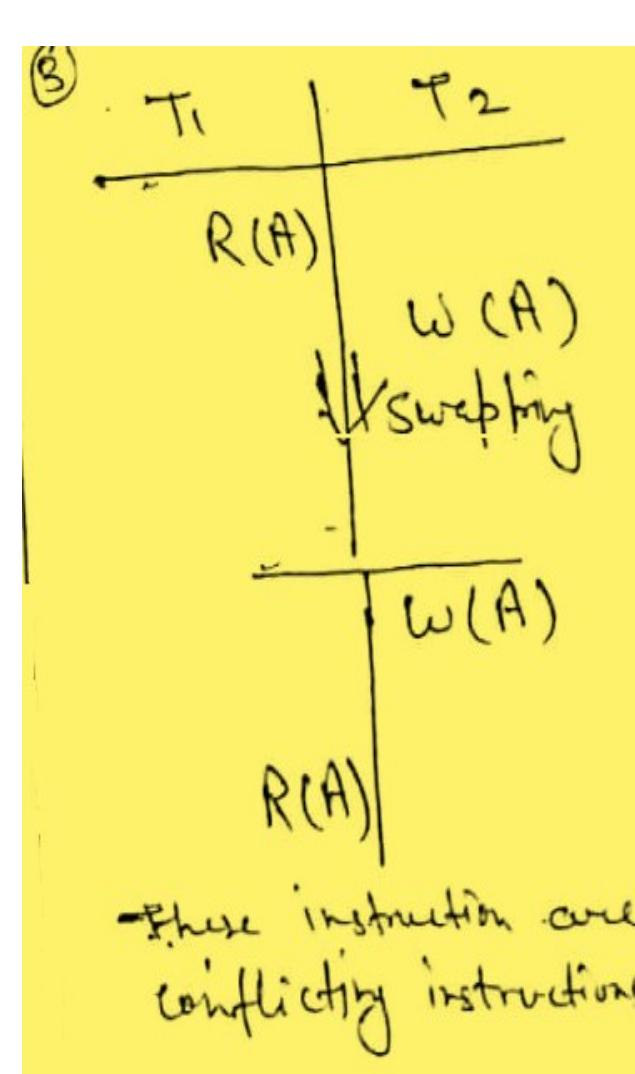
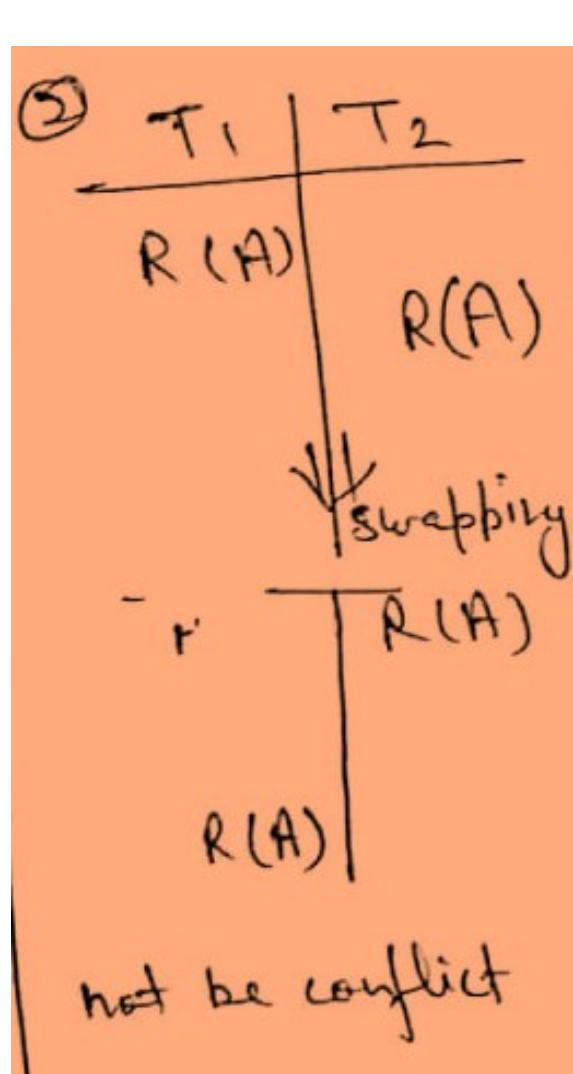
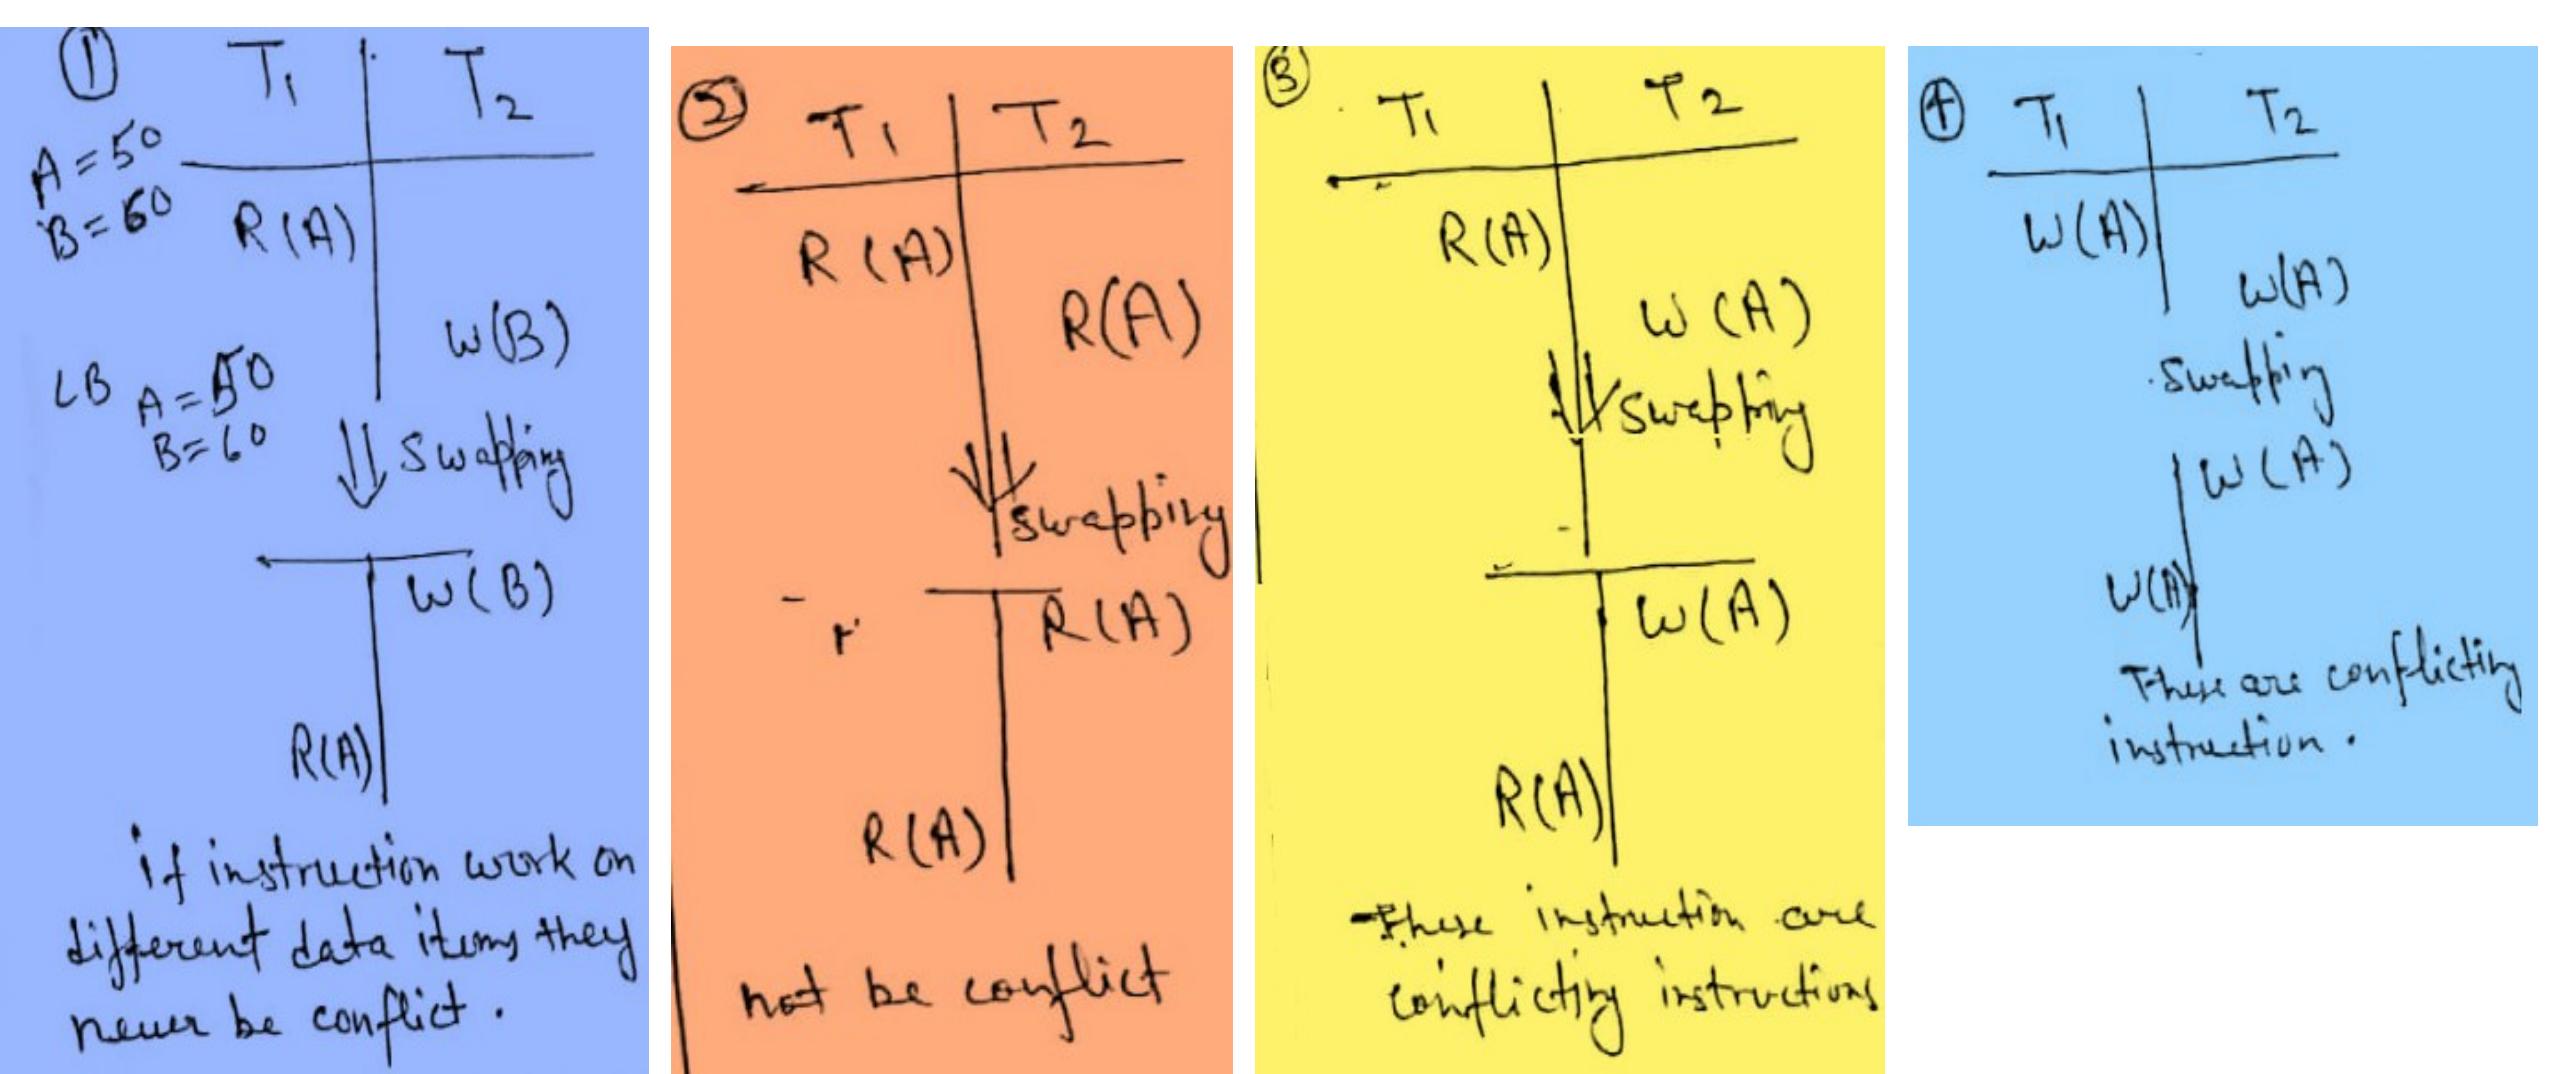
Two operations are called as **conflicting operations** if all the following conditions hold true for them-

- Both the operations belong to different transactions
- Both the operations are on the same data item
- At least one of the two operations is a write operation

Transaction T1	Transaction T2
R1 (A)	
W1 (A)	
	R2 (A)
R1 (B)	

In this schedule,

- W1 (A) and R2 (A) are called as conflicting operations.
- This is because all the above conditions hold true for them.



$S_1$	
$T_1$	$T_2$
$R(A)$	
$w(A)$	$R(A)$
$R(B)$	
$w(B)$	$R(B)$
	$w(B)$

non-Serial

$S_2$	
$T_1$	$T_2$
$R(A)$	
$w(A)$	
$R(B)$	
$w(B)$	
	$R(A)$
	$w(A)$
	$R(B)$
	$w(B)$

Serial

# Checking Whether a Schedule is Conflict Serializable Or Not-

## Step-01:

Find and list all the conflicting operations.

## Step-02:

Start creating a precedence graph by drawing one node for each transaction.

## Step-03:

- Draw an edge for each conflict pair such that if  $X_i$  (V) and  $Y_j$  (V) forms a conflict pair then draw an edge from  $T_i$  to  $T_j$ .
- This ensures that  $T_i$  gets executed before  $T_j$ .

## Step-04:

- Check if there is any cycle formed in the graph.
- If there is no cycle found, then the schedule is conflict serializable otherwise not.

## NOTE-

- By performing the Topological Sort of the Directed Acyclic Graph so obtained, the corresponding serial schedule(s) can be found.
- Such schedules can be more than 1.

# PRACTICE PROBLEMS BASED ON CONFLICT SERIALIZABILITY-

S		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(A)	R(A)	
R(B)	R(B)	
w(A)	w(B)	R(B)

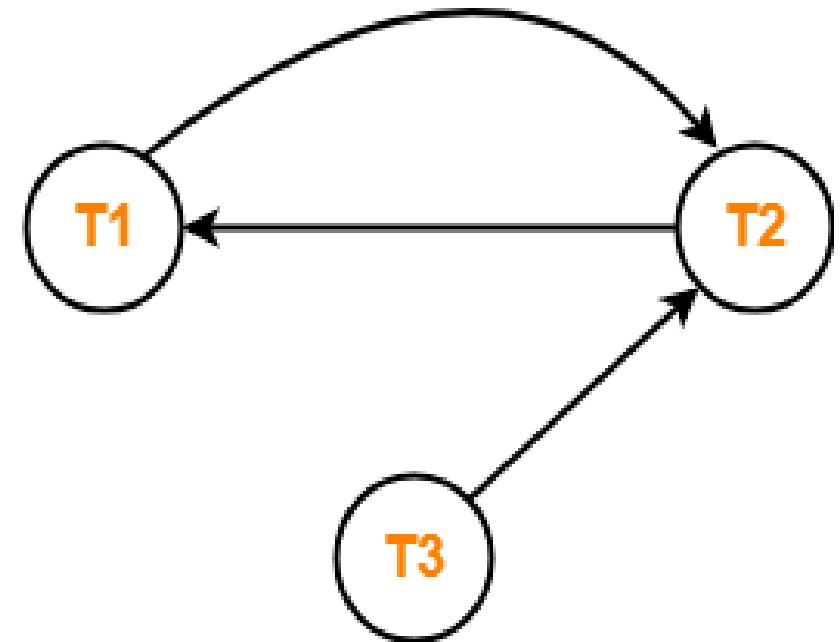
## Step-01:

List all the conflicting operations and determine the dependency between the transactions-

- R<sub>2</sub>(A) , W<sub>1</sub>(A)      (T<sub>2</sub> → T<sub>1</sub>)
- R<sub>1</sub>(B) , W<sub>2</sub>(B)      (T<sub>1</sub> → T<sub>2</sub>)
- R<sub>3</sub>(B) , W<sub>2</sub>(B)      (T<sub>3</sub> → T<sub>2</sub>)

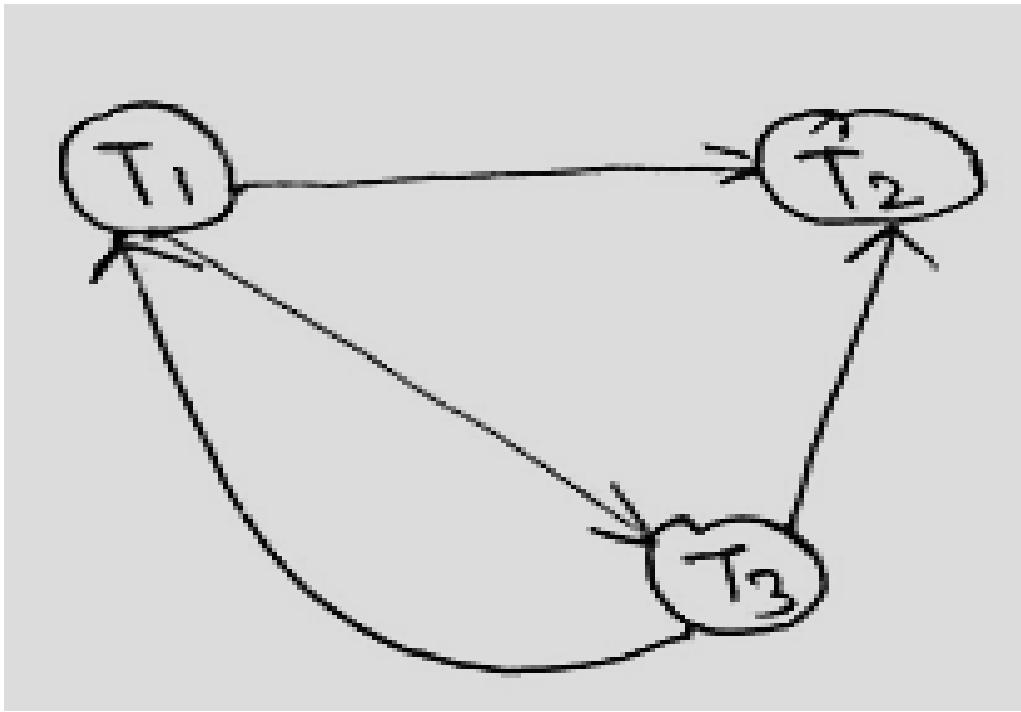
## Step-02:

Draw the precedence graph-



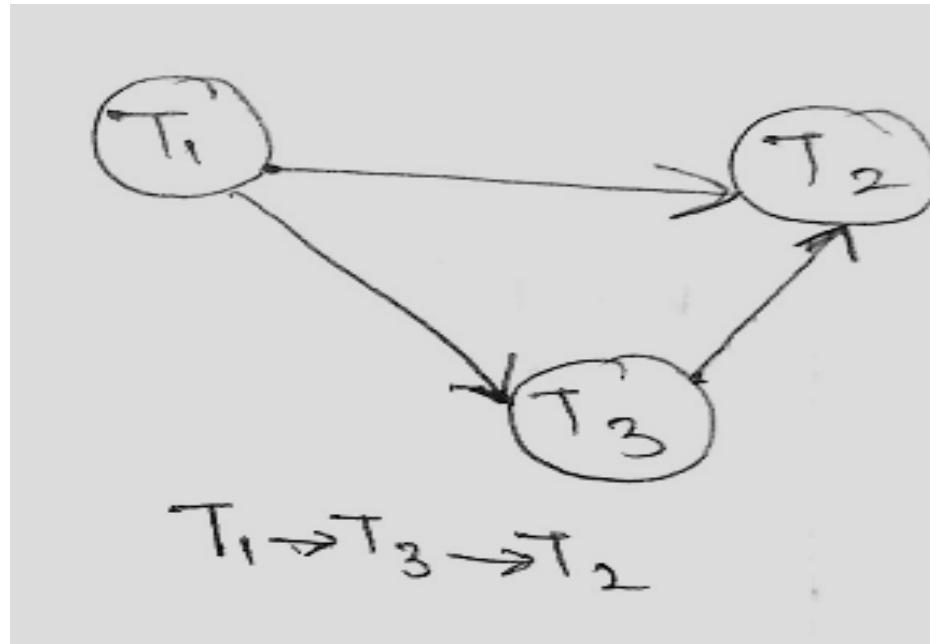
- Clearly, there exists a cycle in the precedence graph.
- Therefore, the given schedule S is not conflict serializable.

$S$		
$T_1$	$T_2$	$T_3$
$R(X)$		
	$R(Y)$	$R(Z)$
$R(Y)$	$w(Y)$	$w(Z)$
		$w(X)$
$w(X)$	$w(Z)$	



- Clearly, there exists a cycle in the precedence graph.
- Therefore, the given schedule  $S$  is not conflict serializable.
- Thus, Number of possible serialized schedules = 0.

S		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(x)		
	R(y)	
w(y)		R(y)
w(x)		
	R(x)	
w(x)		



- Clearly, there exists no cycle in the precedence graph.
- Therefore, the given schedule S is conflict serializable.

possible serialized schedules are-

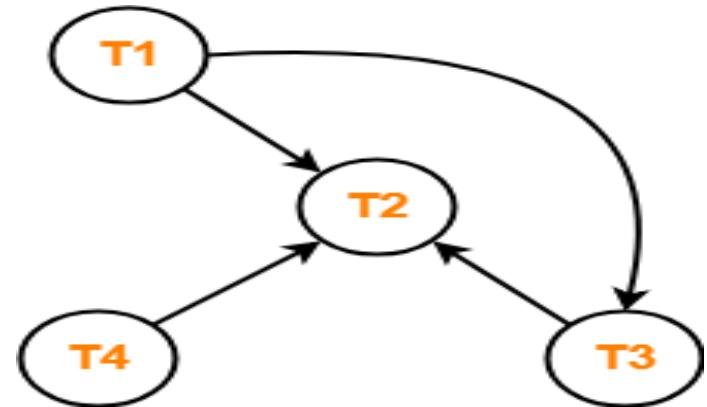
T<sub>1</sub> → T<sub>3</sub> → T<sub>2</sub>

	S		
	$T_1$	$T_2$	$T_3$
$R(a)$		$R(b)$	$R(c)$
			$w(c)$
$w(a)$		$w(b)$	

T1	T2	T3	T4
	R(A)		R(A)
W(B)		R(A)	
	W(A)		
		R(B)	
	W(B)		

conflicting operations and determine the dependency between the transactions-

- $R_4(A), W_2(A)$        $(T_4 \rightarrow T_2)$
- $R_3(A), W_2(A)$        $(T_3 \rightarrow T_2)$
- $W_1(B), R_3(B)$        $(T_1 \rightarrow T_3)$
- $W_1(B), W_2(B)$        $(T_1 \rightarrow T_2)$
- $R_3(B), W_2(B)$        $(T_3 \rightarrow T_2)$
- Clearly, there exists no cycle in the precedence graph.
- Therefore, the given schedule S is conflict serializable



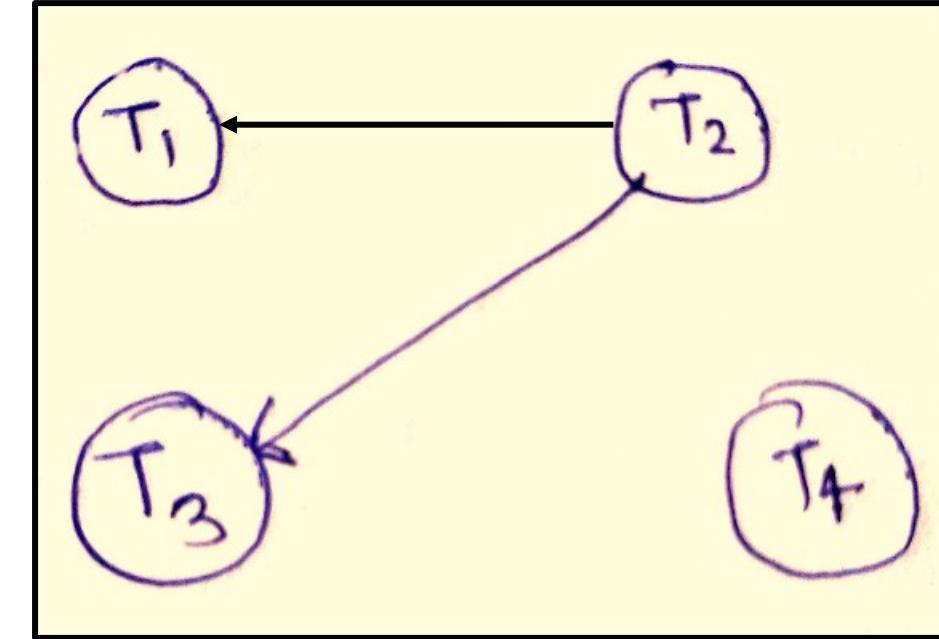
## Finding the Serialized Schedules-

- All the possible topological orderings of the above precedence graph will be the possible serialized schedules.
- The topological orderings can be found by performing the Topological Sort of the above precedence graph.

After performing the topological sort, the possible serialized schedules are-

- $T_1 \rightarrow T_3 \rightarrow T_4 \rightarrow T_2$
- $T_1 \rightarrow T_4 \rightarrow T_3 \rightarrow T_2$
- $T_4 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2$

T1	T2	T3	T4
	R(X)		
W(X) Commit		W(X) Commit	
	W(Y) R(Z) Commit		
			R(X) R(Y) Commit



## Checking Whether S is Recoverable Or Not-

- . Conflict serializable schedules are always recoverable.
- Therefore, the given schedule S is recoverable.

**Alternatively,**

- There exists no dirty read operation.
- This is because all the transactions which update the values commits immediately.

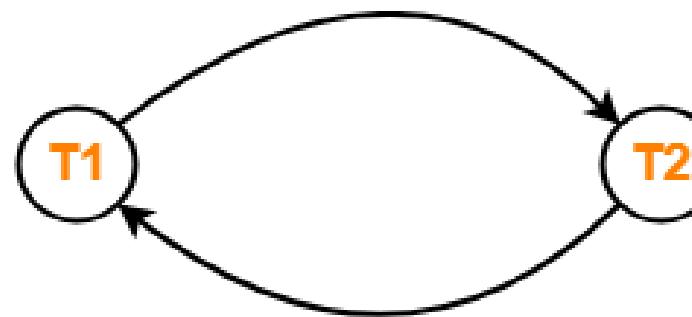
- Clearly, there exists no cycle in the precedence graph.
- Therefore, the given schedule S is conflict serializable.

T1	T2
R(A)	
A = A-10	
	R(A)
	Temp = 0.2 x A
	W(A)
	R(B)
W(A)	
R(B)	
B = B+10	
W(B)	
	B = B+Temp
	W(B)

T1	T2
R(A)	
	R(A)
	W(A)
	R(B)
W(A)	
R(B)	
	W(B)
	W(B)

conflicting operations and determine the dependency between the transactions-

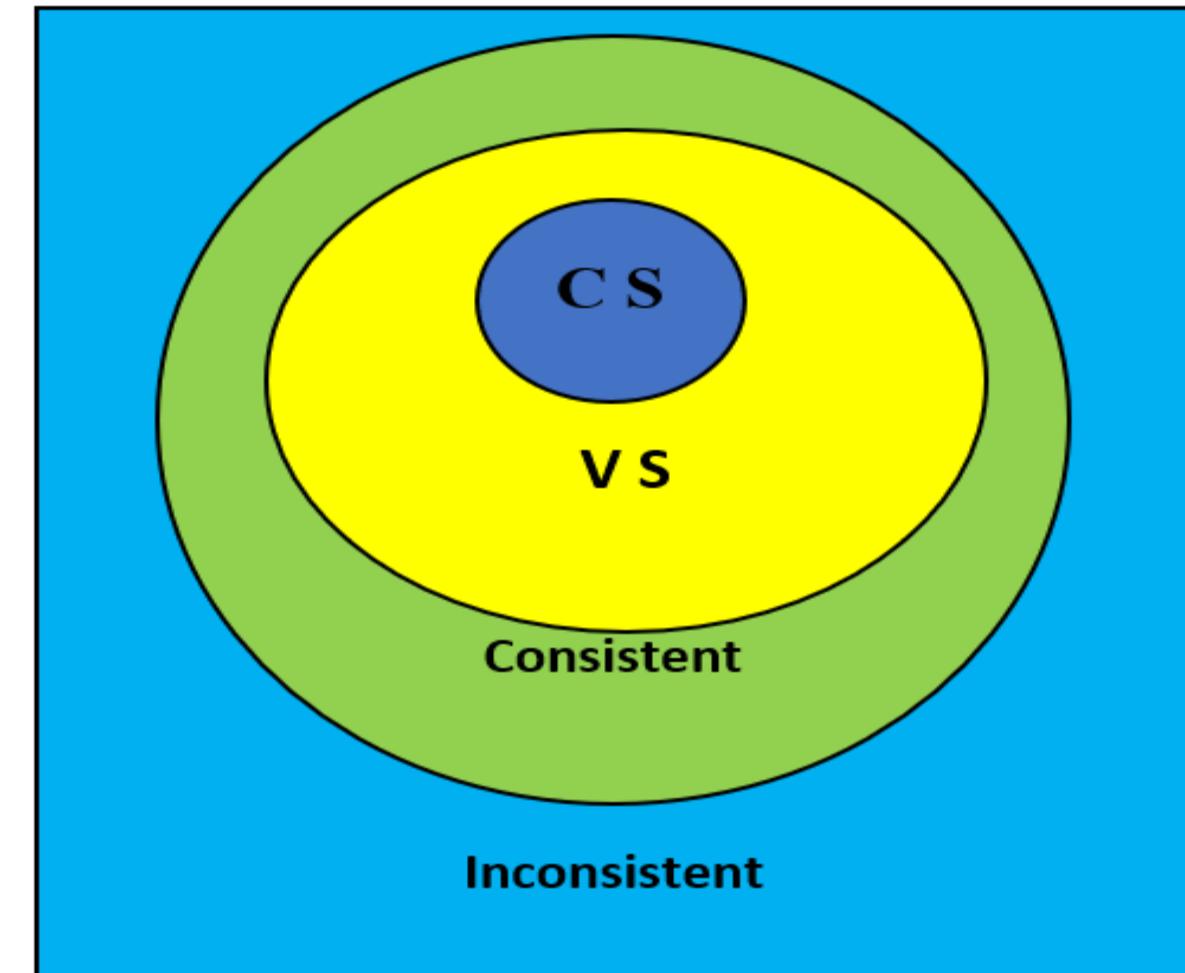
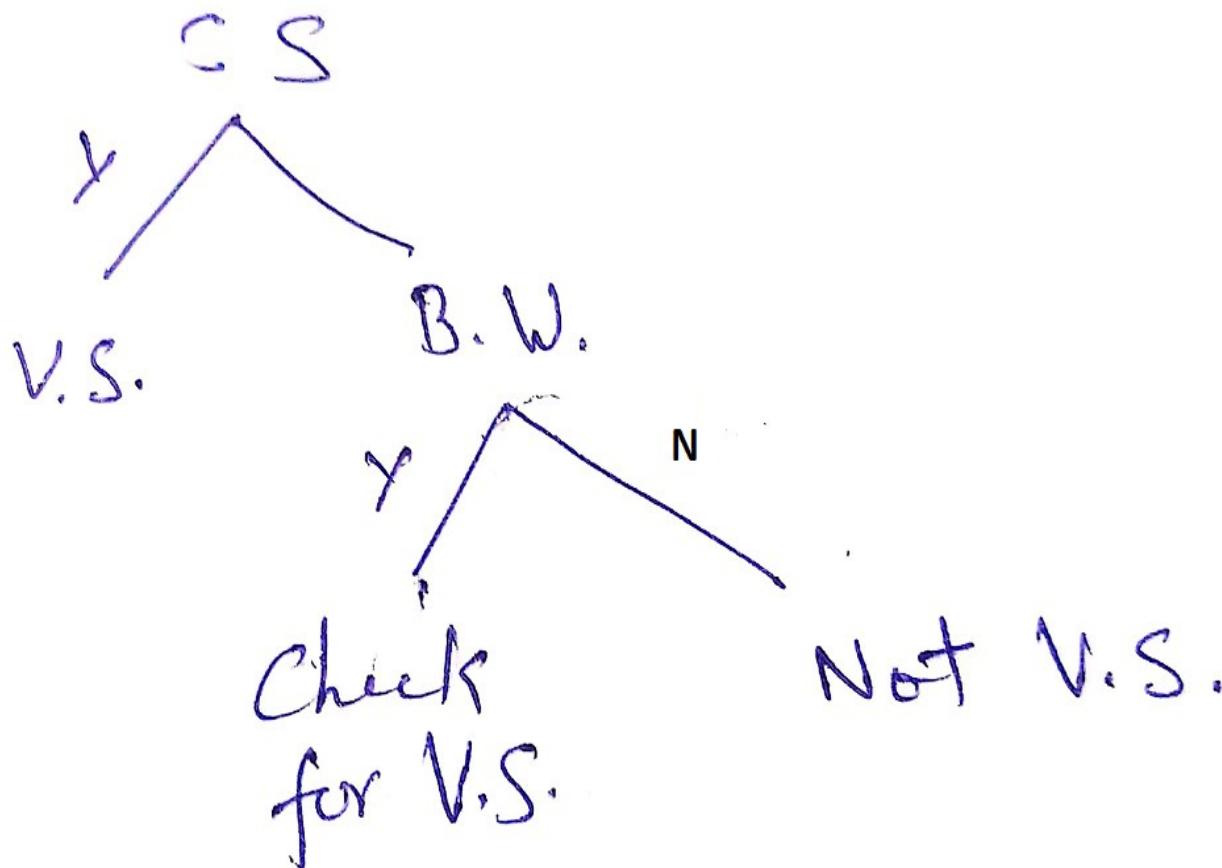
- $R_1(A), W_2(A)$   $(T_1 \rightarrow T_2)$
- $R_2(A), W_1(A)$   $(T_2 \rightarrow T_1)$
- $W_2(A), W_1(A)$   $(T_2 \rightarrow T_1)$
- $R_2(B), W_1(B)$   $(T_2 \rightarrow T_1)$
- $R_1(B), W_2(B)$   $(T_1 \rightarrow T_2)$
- $W_1(B), W_2(B)$   $(T_1 \rightarrow T_2)$



- Clearly, there exists a cycle in the precedence graph.
- Therefore, the given schedule S is not conflict serializable.
- Thus, Number of possible serialized schedules = 0.

# View Serializability in DBMS

- A schedule is view serializable when it is view equivalent to a serial schedule.
- All conflict serializable schedules are view serializable.
- The view serializable which is not a conflict serializable contains blind writes.



To check whether a given schedule is view serializable, we need to check whether the given schedule is **View Equivalent** to its serial schedule

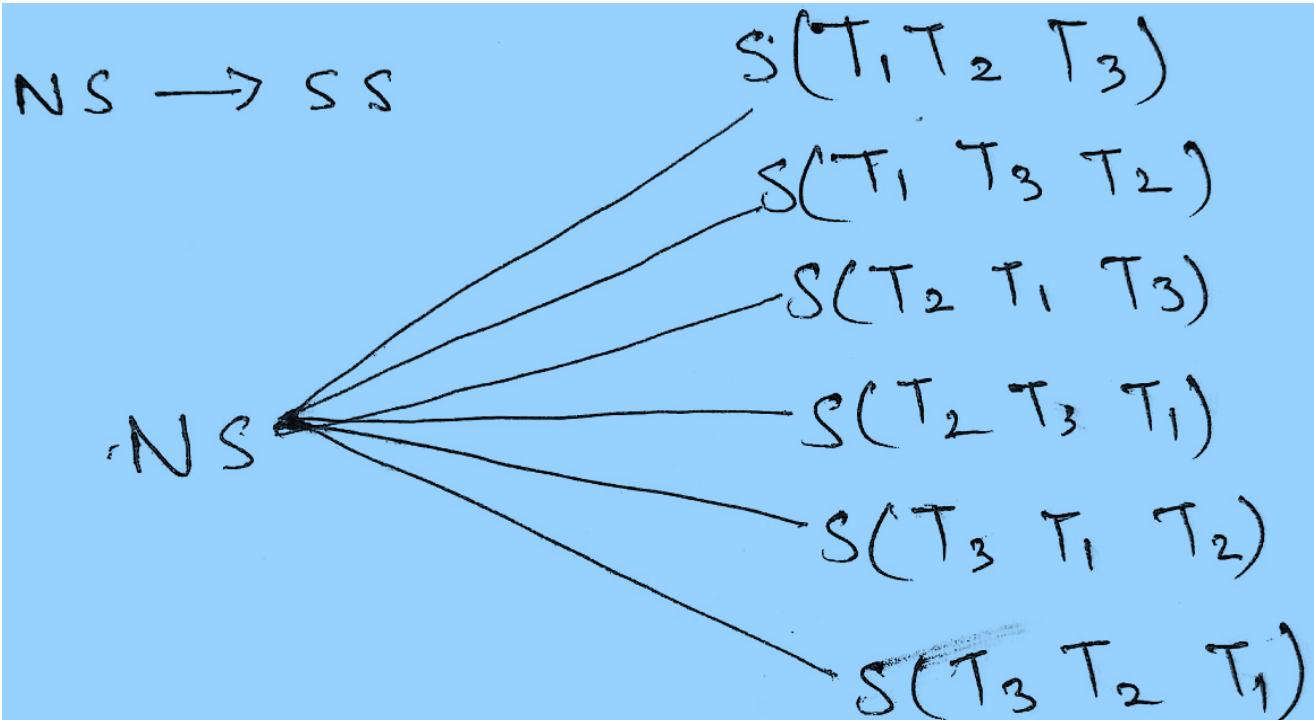
S	T1	T2
R(X)		
W(X)		
	R(X)	
	W(X)	
R(Y)		
W(Y)		
	R(Y)	
	W(Y)	

S	T1	T2
R(X)		
W(X)		
	R(Y)	
	W(Y)	
R(X)		
W(X)		
	R(Y)	
	W(Y)	

If we can prove that the given schedule is **View Equivalent** to its serial schedule then the given schedule is called **view Serializable**.

## Why we need View Serializability?

We know that a serial schedule never leaves the database in inconsistent state because there are no concurrent transactions execution. However a non-serial schedule can leave the database in inconsistent state because there are multiple transactions running concurrently. By checking that a given non-serial schedule is view serializable, we make sure that it is a consistent schedule.



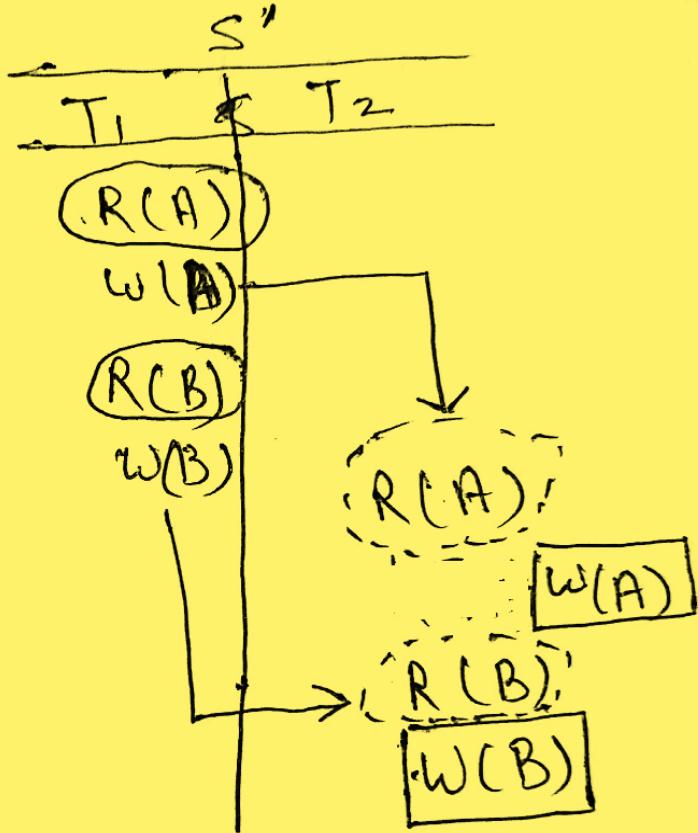
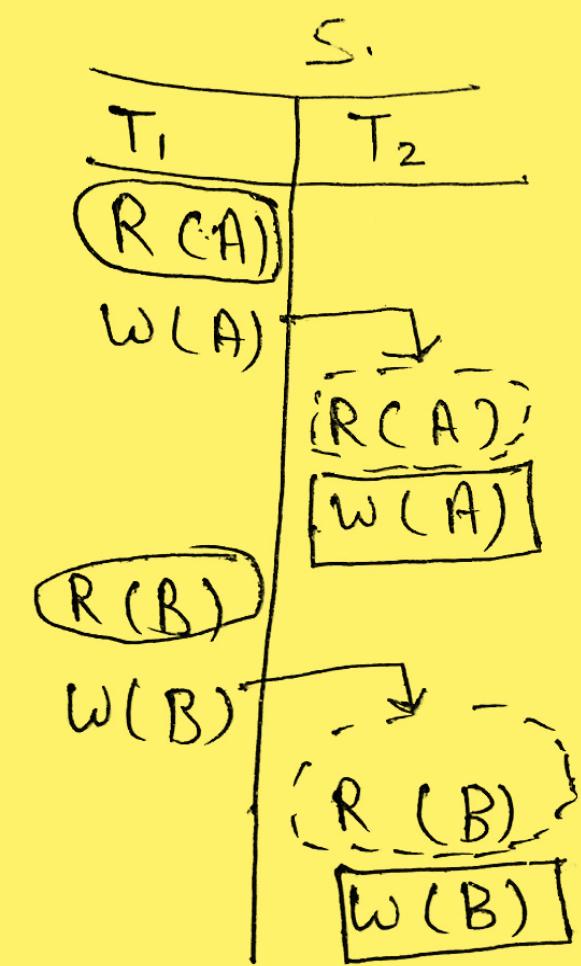
# View Equivalent

**1. Initial Read:** Initial read of each data item in transactions must match in both schedules. For example, if transaction T1 reads a data item X before transaction T2 in schedule S1 then in schedule S2, T1 should read X before T2.

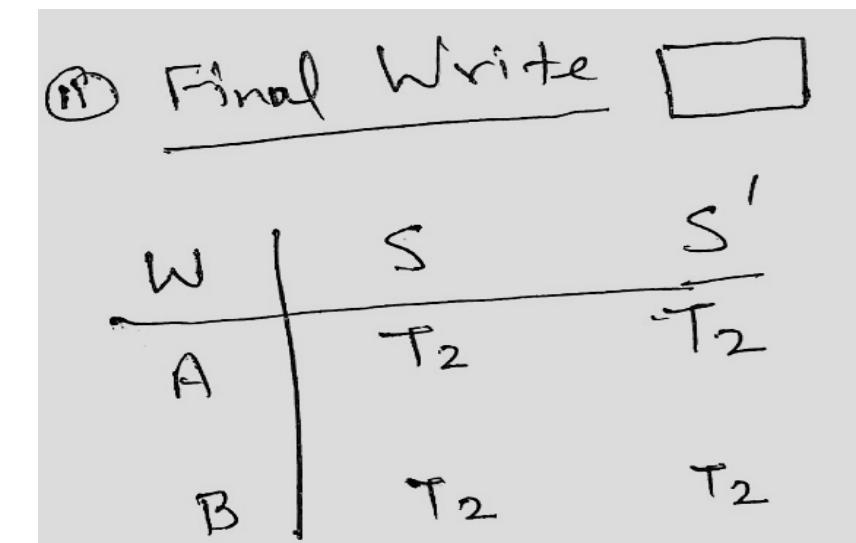
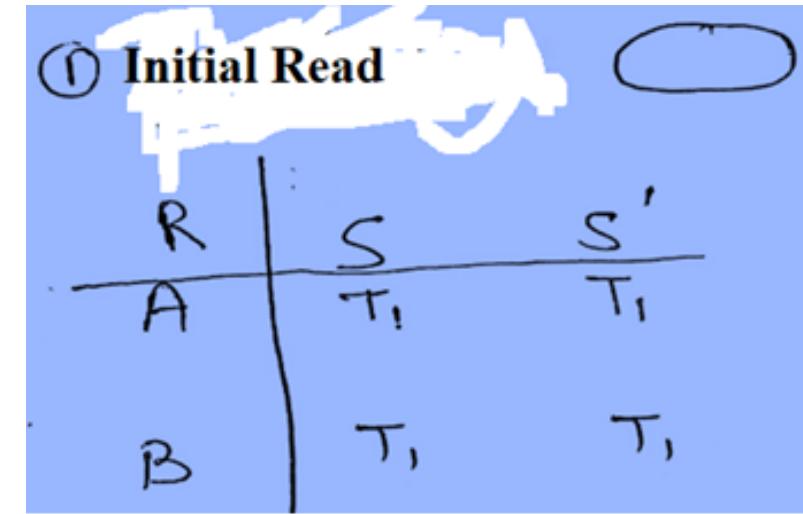
**Read vs Initial Read:** You may be confused by the term initial read. Here initial read means the first read operation on a data item, for example, a data item X can be read multiple times in a schedule but the first read operation on X is called the initial read. This will be more clear once we will get to the example in the next section of this same article.

**2. Final Write:** Final write operations on each data item must match in both the schedules. For example, a data item X is last written by Transaction T1 in schedule S1 then in S2, the last write operation on X should be performed by the transaction T1.

**3. Update Read:** If in schedule S1, the transaction T1 is reading a data item updated by T2 then in schedule S2, T1 should read the value after the write operation of T2 on same data item. For example, In schedule S1, T1 performs a read operation on X after the write operation on X by T2 then in S2, T1 should read the X after T2 performs write on X.



(III) Update Read / Intermediate Read



All the three conditions that checks whether the two schedules are view equivalent are satisfied which means  $S_1$  and  $S_2$  are view equivalent. So we can say that the schedule  $S_1$  is view serializable schedule.

## Non-Serial

S1	
T1	T2
R(X)	
W(X)	
	R(X)
	W(X)
R(Y)	
W(Y)	
	R(Y)
	W(Y)

## Serial

S2	
T1	T2
R(X)	
W(X)	
	R(Y)
	W(Y)
R(X)	
W(X)	
	R(Y)
	W(Y)

## Initial Read

In schedule S1, transaction T1 first reads the data item X.

In S2 also transaction T1 first reads the data item X.

Lets check for Y. In schedule S1, transaction T1 first reads the data item Y. In S2 also the first read operation on Y is performed by T1.

## Final Write

In schedule S1, the final write operation on X is done by transaction T2. In S2 also transaction T2 performs the final write on X.

Lets check for Y. In schedule S1, the final write operation on Y is done by transaction T2. In schedule S2, final write on Y is done by T2.

## Update Read

In S1, transaction T2 reads the value of X, written by T1. In S2, the same transaction T2 reads the X after it is written by T1.

In S1, transaction T2 reads the value of Y, written by T1. In S2, the same transaction T2 reads the value of Y after it is updated by T1.

Since all the three conditions that checks whether the two schedules are view equivalent are satisfied in this example, which means S1 and S2 are view equivalent. Also, as we know that the schedule S2 is the serial schedule of S1, thus we can say that the schedule S1 is view serializable schedule.

Q. Check given schedule is View Serializable or not

S		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(a)		
	w(a)	
w(a)		w(a)

S <sub>1</sub>		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(a)		
w(a)	w(a)	
	w(a)	w(a)

S<sub>1</sub> → T<sub>1</sub> T<sub>2</sub> T<sub>3</sub> ✓  
S<sub>2</sub> → T<sub>1</sub> T<sub>3</sub> T<sub>2</sub>  
S<sub>3</sub> → T<sub>2</sub> T<sub>1</sub> T<sub>3</sub>  
S<sub>4</sub> → T<sub>2</sub> T<sub>3</sub> T<sub>1</sub>  
S<sub>5</sub> → T<sub>3</sub> T<sub>1</sub> T<sub>2</sub>  
S<sub>6</sub> → T<sub>3</sub> T<sub>2</sub> T<sub>1</sub>

- ① Initial Read
- ② Final Write
- ③ Update Read

All the three conditions that checks whether the two schedules are view equivalent are satisfied which means S<sub>1</sub> and S<sub>2</sub> are view equivalent. So we can say that the schedule S<sub>1</sub> is view serializable schedule.

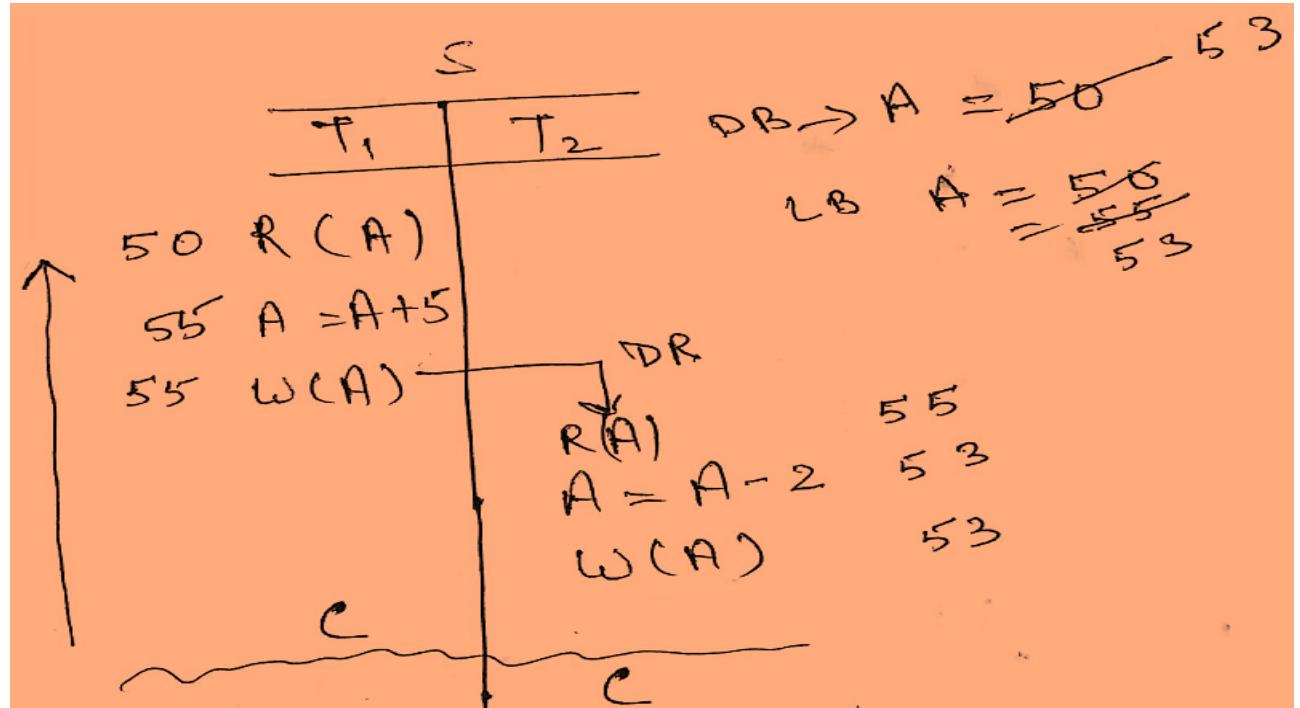
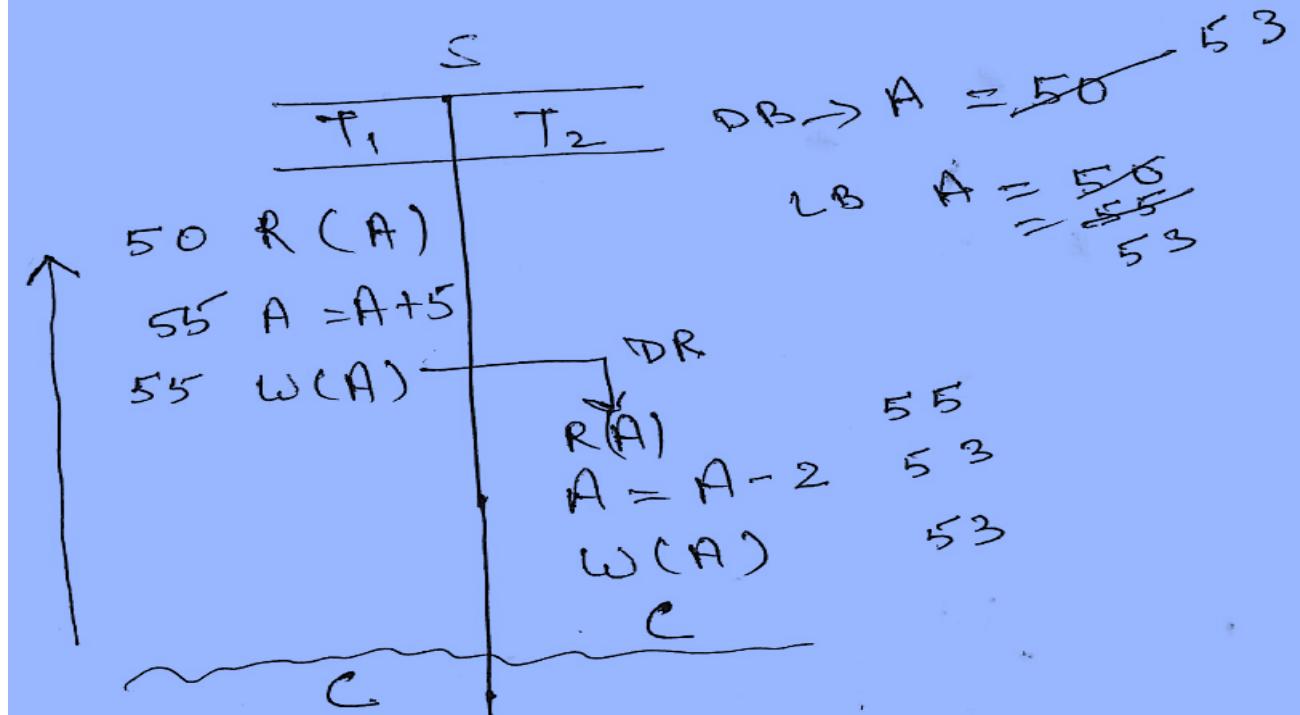
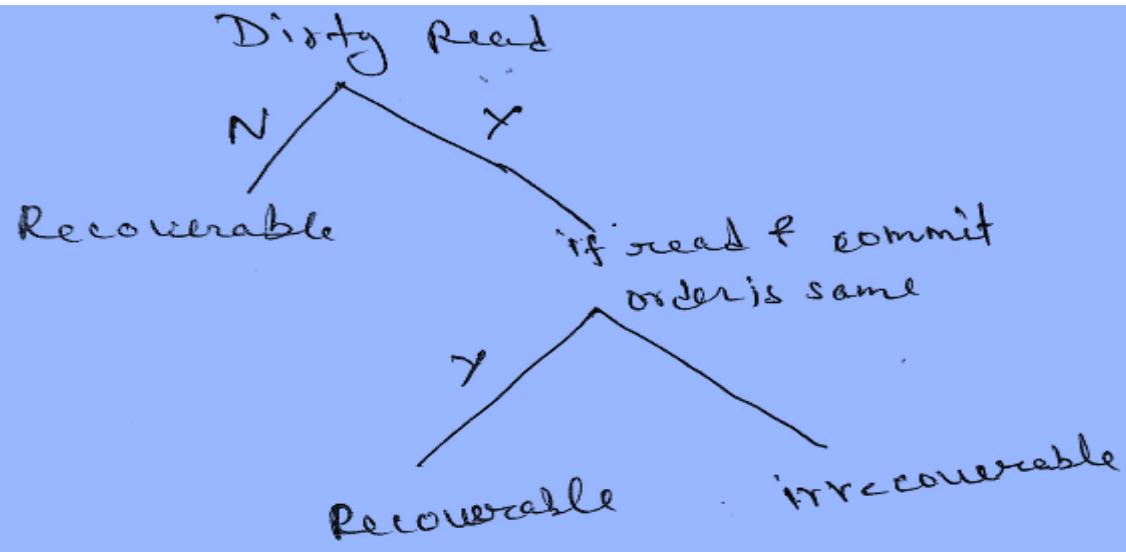
# Recoverability

- Recoverable Schedule
- Irrecoverable Schedules-

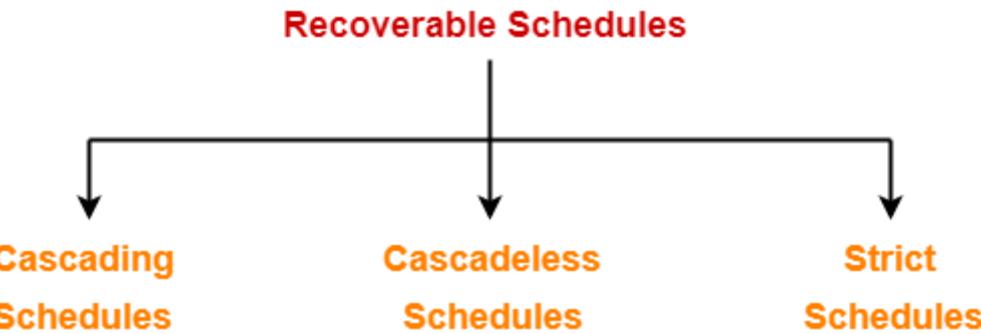
If in a schedule,

- A transaction performs a dirty read operation from an uncommitted transaction
- And commits before the transaction from which it has read the value

then such a schedule is known as  
**an Irrecoverable Schedule.**

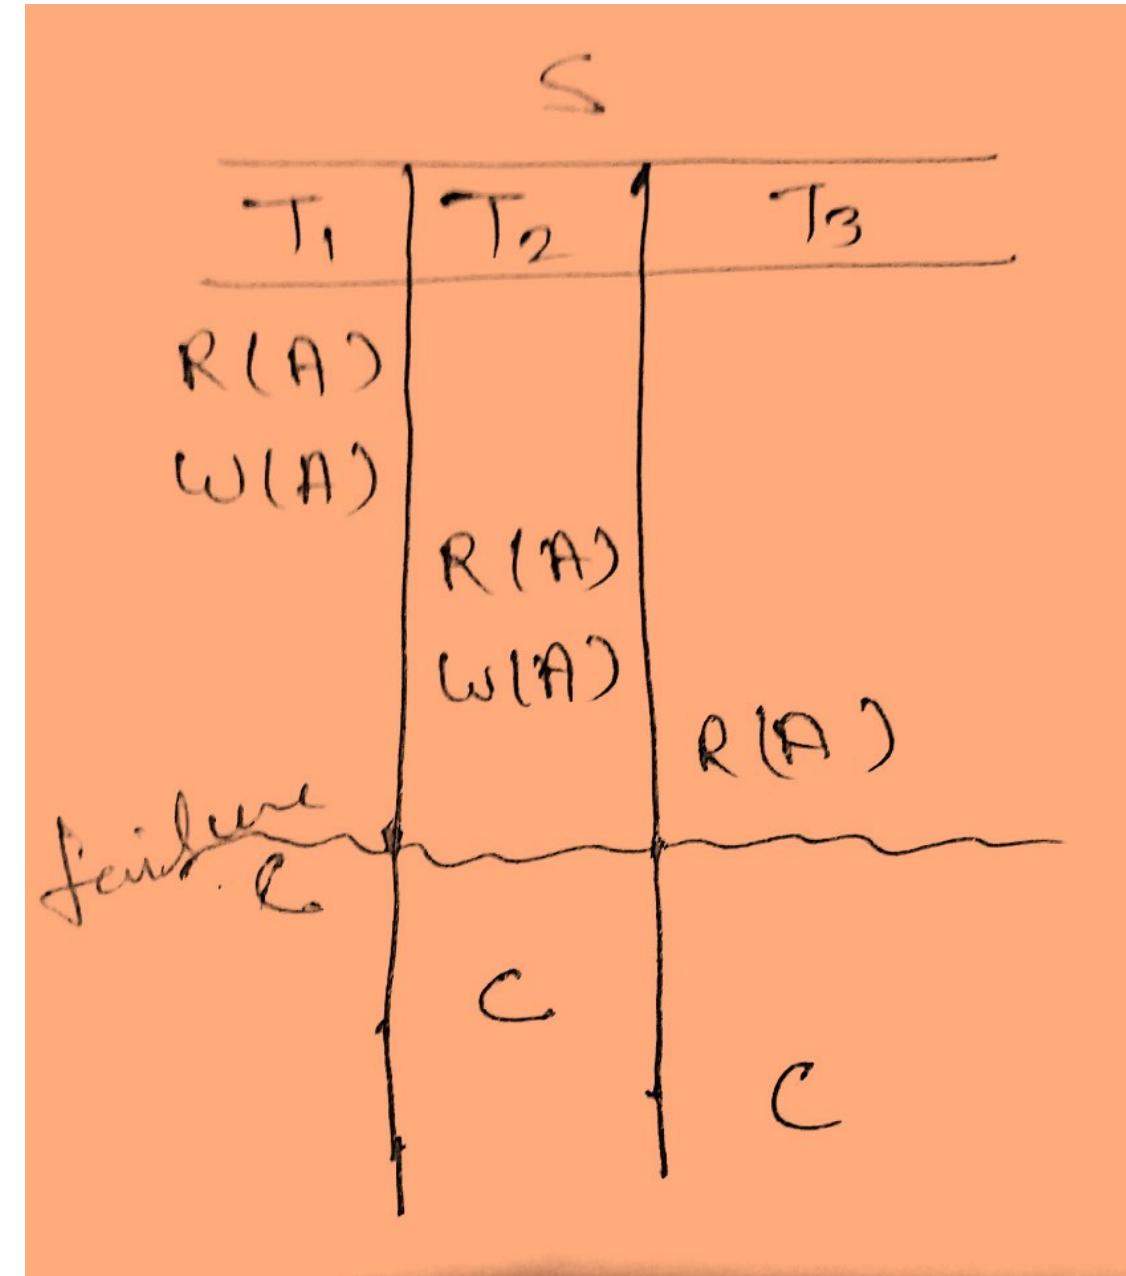
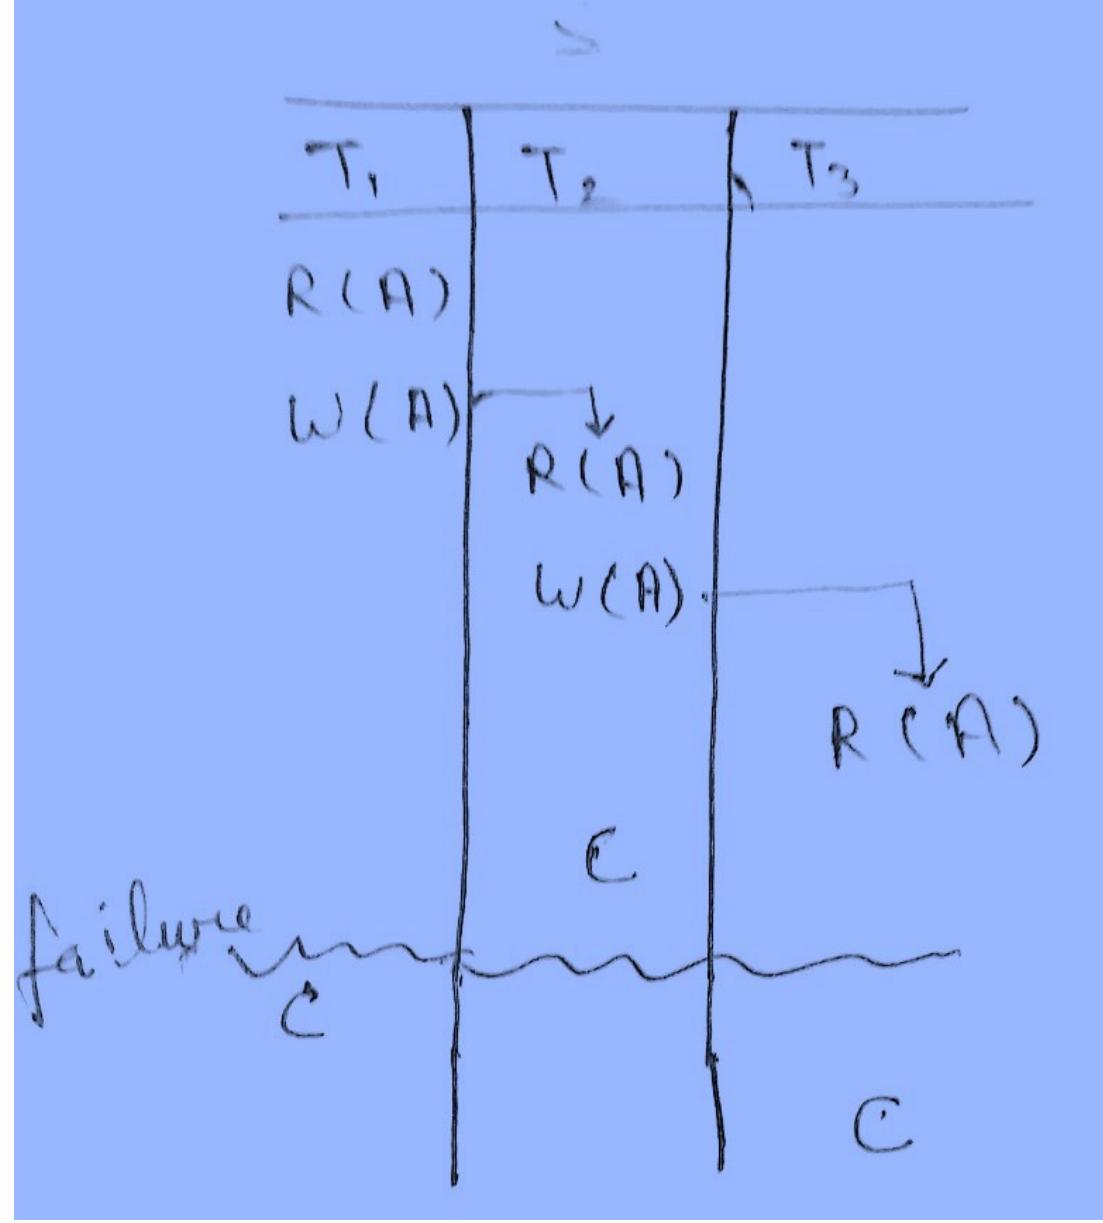


## Types of Recoverable Schedules-



## Cascading Schedule-

- If in a schedule, failure of one transaction causes several other dependent transactions to rollback or abort, then such a schedule is called as a **Cascading Schedule** or **Cascading Rollback** or **Cascading Abort**.
- It simply leads to the wastage of CPU time.



# Cascadeless Schedule-

If in a schedule, a transaction is not allowed to read a data item until the last transaction that has written it is committed or aborted, then such a schedule is called as a **Cascadeless Schedule**.

In other words,

- Cascadeless schedule allows only committed read operations.
- Therefore, it avoids cascading roll back and thus saves CPU time.

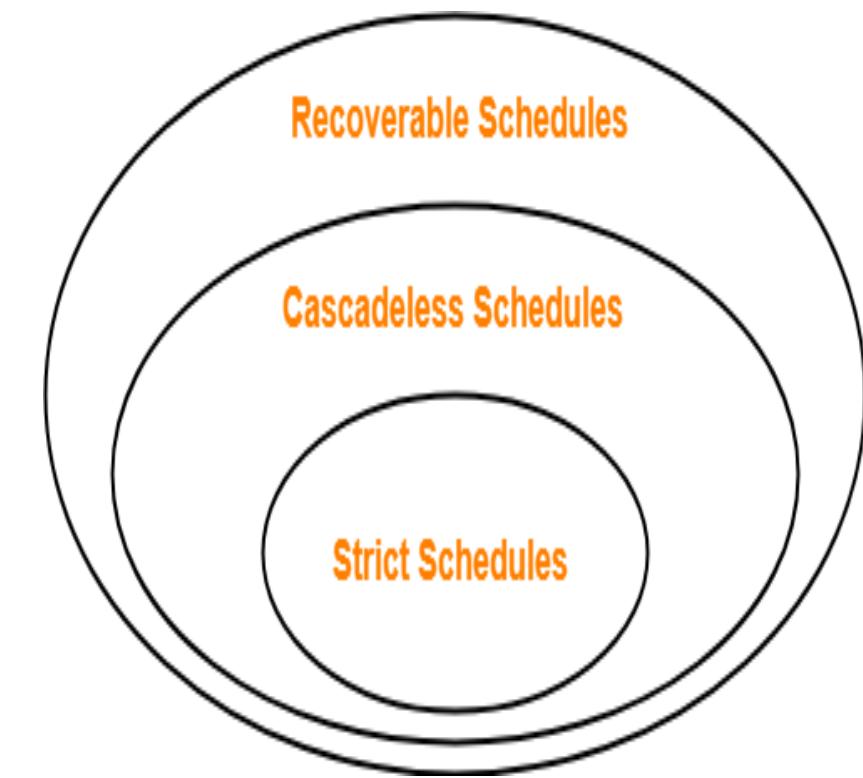
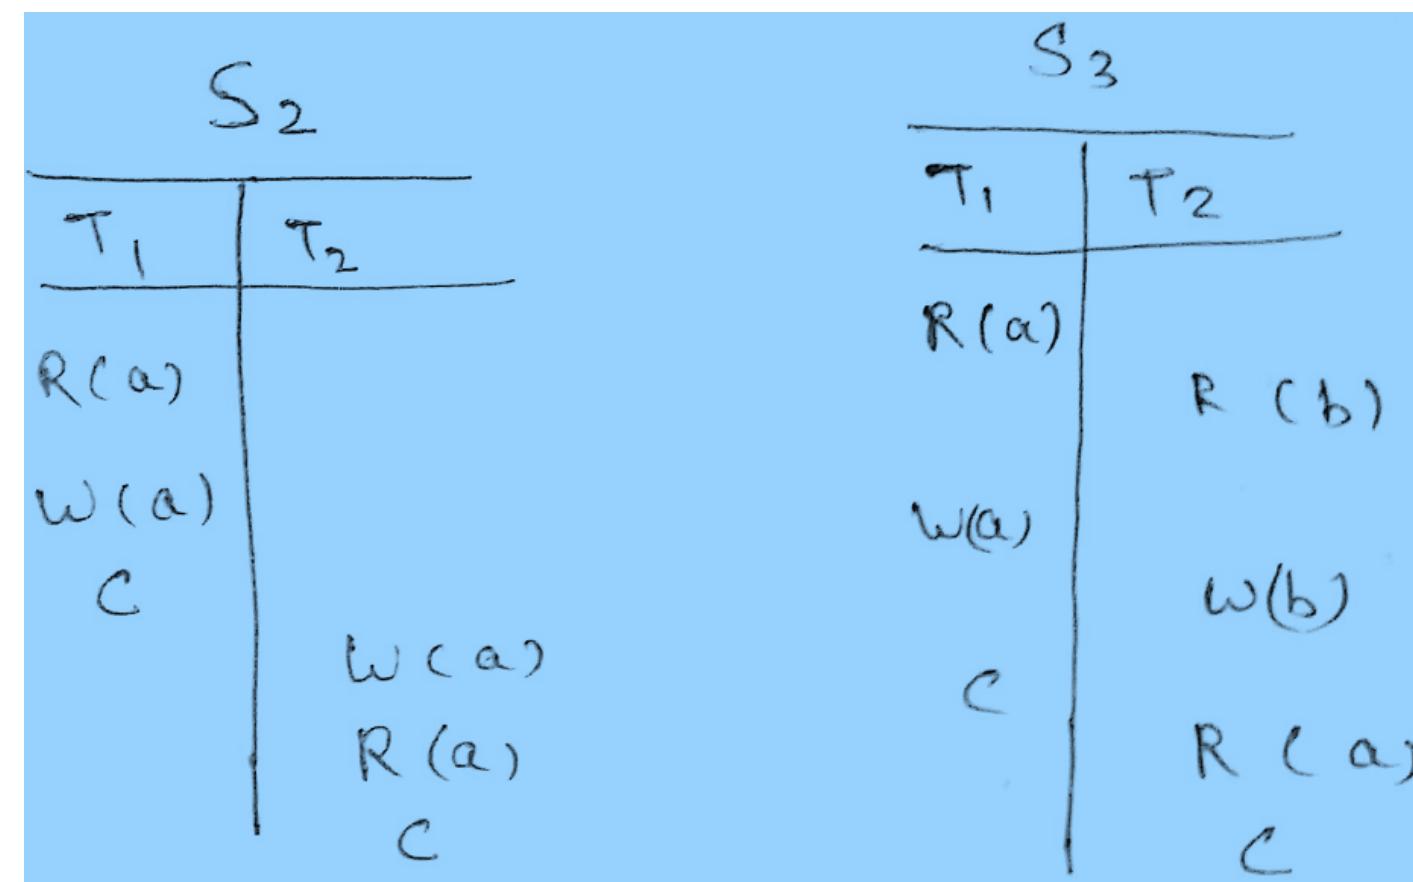
S		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(A)		
W(A)		
	R(A)	
	W(A)	
failure		
	C	
		C

S		
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
R(A)		
W(A)		
C		
	R(A)	
	W(A)	
	C	
		R(A)
		C

# Strict Schedule-

If in a schedule, a transaction is neither allowed to read nor write a data item until the last transaction that has written it is committed or aborted, then such a schedule is called as a **Strict Schedule**.

- Strict schedule allows only committed read and write operations.
- Clearly, strict schedule implements more restrictions than cascadeless schedule



# CONCURRENCY CONTROL

- Now we understand that if there is a schedule, we can check whether it will work correctly or not, i.e., whether it will maintain the consistency of the database or not. (conflict serializability, view serializability, recoverability, and cascade less)
- Now we will understand those protocols that guarantee how to design those schedules that ensure conflict serializability or other properties. There are different approaches or ideas to ensure conflict serializability, which is the most important property.
- So first we must understand what is the possibility of conflict is between two instructions, and if somehow we manage, then the generated schedule will always be conflict serializable.

## Goals of a Protocol: -

We desire the following properties from the schedule-generating protocols

- Concurrency should be as high as possible, as this is our ultimate goal, for which we are making all the effort.
- The time taken by a transaction should also be less.
- Desirable Properties satisfied by the protocol
- Easy to understand and implement

**Time stamping based method:** - Where, before entering the system, a specific order is decided among the transactions, so in case of a clash, we can decide which one to allow and which to stop.

• **Lock-based method:** - where we ask a transaction to first lock a data item before using it. So that no different transaction can use a data at the same time, removing any possibility of conflict.

## 2-phase locking

1. Basic 2pl
2. Conservative 2pl
3. Rigorous 2pl
4. Strict 2pl

• **Graph-based protocol**

• **Validation-based protocol** – The Majority of transactions are read-only transactions, the rate of conflicts among transactions may be low, thus many transactions, if executed without the supervision of a concurrency control scheme, would nevertheless leave the system in a consistent state.

# TIME STAMP ORDERING PROTOCOL

- The basic idea of time stamping is to decide the order between transactions before they enter the system using a stamp (time stamp). In case of any conflict during execution, the order can be decided using the time stamp.
- Let's understand how this protocol works. Here, we have two ideas of timestamping, one for the transaction and the other for the data item.

## Time stamp for transaction

- With each transaction  $t_i$ , in the system, we associate a unique fixed timestamp, denoted by  $TS(t_i)$ .
- This timestamp is assigned by the database system to a transaction at the time the transaction enters the system.
- If a transaction has been assigned a timestamp  $TS(t_i)$  and a new transaction  $t_j$  , enters into the system with a timestamp  $TS(t_j)$ , then always  $TS(t_i) < TS(t_j)$ .

### \*\*\*Note

1. First time stamp of a transaction remain fixed throughout the execution
2. Second it is unique means no two transaction can have the same timestamp.

Time stamp with data item, in order to assure such scheme, the protocol maintains for each data item Q two timestamp values:

1. **W-timestamp(Q)** is the largest time-stamp of any transaction that successfully executed write(Q).
2. **R-timestamp(Q)** is the largest time-stamp of any transaction that successfully executed read(Q).

These timestamps are updated whenever a new read(Q) or write(Q) instruction is executed.

### Suppose a transaction Ti request a *read(Q)*

1. If  $TS(Ti) < W\text{-timestamp}(Q)$ , then  $Ti$  needs to read a value of  $Q$  that was already overwritten. Hence, the read operation is rejected, and  $Ti$  is rolled back.
2. If  $TS(Ti) \geq W\text{-timestamp}(Q)$ , then the read operation is executed, and  $R\text{-timestamp}(Q)$  is set to the maximum of  $R\text{-timestamp}(Q)$  and  $TS(Ti)$ .

Suppose that transaction  $Ti$  issues  $\text{write}(Q)$ .

1. If  $TS(Ti) < R\text{-timestamp}(Q)$ , then the value of  $Q$  that  $Ti$  is producing was needed previously, and the system assumed that that value would never be produced. Hence, the write operation is rejected, and  $Ti$  is rolled back.
2. If  $TS(Ti) < W\text{-timestamp}(Q)$ , then  $Ti$  is attempting to write an obsolete value of  $Q$ . Hence, this write operation is rejected, and  $Ti$  is rolled back.
3. If  $TS(Ti) \geq R\text{-timestamp}(Q)$ , then the write operation is executed, and  $W\text{-timestamp}(Q)$  is set to  $\max(W\text{-timestamp}(Q), TS(Ti))$ .
4. If  $TS(Ti) \geq W\text{-timestamp}(Q)$ , then the write operation is executed, and  $W\text{-timestamp}(Q)$  is set to  $\max(W\text{-timestamp}(Q), TS(Ti))$ .

	Conflict Serializability	View Serializability	Recoverability	Cascadelessness	Deadlock Freedom
Time Stamp Ordering	YES	YES	NO	NO	YES
Thomas Write Rule					
Basic 2PL					
Conservative 2PL					
Rigorous 2PL					
Strict 2PL					

## THOMAS WRITE RULE

- Thomas write is an improvement in time stamping protocol, which makes some modification and may generate those protocols that are even view serializable, because it allows greater potential concurrency.
- It is a Modified version of the timestamp-ordering protocol in which Blind write operations may be ignored under certain circumstances.
- The protocol rules for read operations remain unchanged. while for write operation, there is slightly change in Thomas write rule than timestamp ordering protocol.

**When  $T_i$  attempts to write data item  $Q$ ,**

- if  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $\{Q\}$ . Rather than rolling back  $T_i$  as the timestamp ordering protocol would have done, this {write} operation can be ignored.

- This modification is valid as the any transaction with  $TS(T_i) < W\text{-timestamp}(Q)$ , the value written by this transaction will never be read by any other transaction performing  $\text{Read}(Q)$  ignoring such obsolete write operation is considerable.
- Thomas' Write Rule allows greater potential concurrency. Allows some view-serializable schedules that are not conflict serializable.

$T_3$	$T_4$	$T_6$
$\text{read}(Q)$	$\text{write}(Q)$	
$\text{write}(Q)$		$\text{write}(Q)$

	Conflict Serializability	View Serializability	Recoverability	Cascadelessness	Deadlock Freedom
Time Stamp Ordering	YES	YES	NO	NO	YES
Thomas Write Rule	NO	YES	NO	NO	YES
Basic 2PL					
Conservative 2PL					
Rigorous 2PL					
Strict 2PL					

## Lock Based Protocols

- ❑ To ensure isolation, it is required that data items be accessed in a mutually exclusive manner, i.e., while one transaction is accessing a data item, no other transaction can modify that data item. Locking is the most fundamental approach to ensure this.
- ❑ Lock-based protocols ensure this requirement. Idea is first obtain a lock on the desired data item then if lock is granted then perform the operation and then unlock it.

In general, we support two modes of lock because, to provide better concurrency.

### ***Shared mode***

If transaction  $T_i$  has obtained a shared-mode lock (denoted by  $S$ ) on any data item  $Q$ , then  $T_i$  can read, but cannot write  $Q$ , any other transaction can also acquire a shared mode lock on the same data item(this is the reason we called this shared mode).

### ***Exclusive mode***

If transaction  $T_i$  has obtained an exclusive-mode lock (denoted by  $X$ ) on any data item  $Q$ , then  $T_i$  can both read and write  $Q$ , any other transaction cannot acquire either a shared or exclusive mode lock on the same data item. (this is the reason we called this exclusive mode)

## Lock –Compatibility Matrix

Shared is compatible only with shared, while exclusive is not compatible with either shared or exclusive.

To access a data item, transaction  $T_i$  must first lock that item, if the data item is already locked by another transaction in an incompatible mode, or some other transaction is already waiting in non-compatible mode, then concurrency control manager will not grant the lock until all incompatible locks held by other transactions have been released. The lock is then granted.

		Current State of lock of data items		
		Exclusive	Shared	Unlocked
Requested Lock	Exclusive	N	N	Y
	Shared	N	Y	Y
	Unlock	Y	Y	-

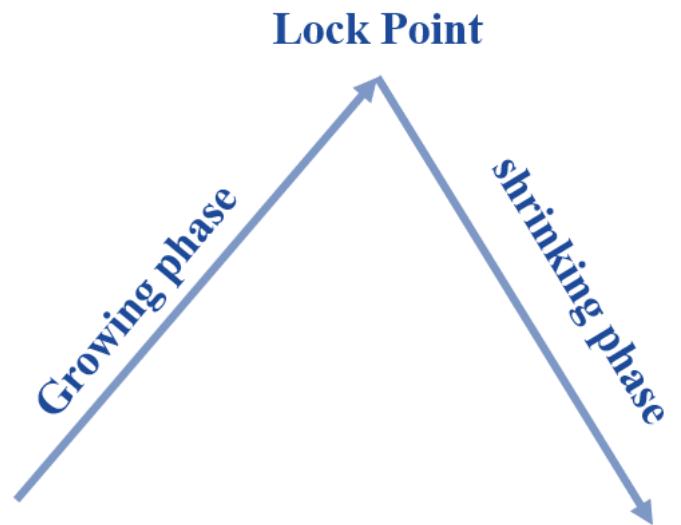
Lock based protocol *do not ensure serializability* as granting and releasing of lock do not follow any order and any transaction any time may go for lock and unlock. Here in the example below we can see, that even this transaction in using locking but neither it is conflict serializable nor independent from deadlock.

- If we do not use locking, or if we unlock data items too soon after reading or writing them, we may get inconsistent states, as there exists a possibility of dirty read. On the other hand, if we do not unlock a data item before requesting a lock on another data item, concurrency will be poor.
- We shall require that each transaction in the system follow a set of rules, called a **locking protocol**, indicating when a transaction may lock and unlock each of the data items for e.g. 2pl or graph based locking.
- Locking protocols restrict the number of possible schedules.

$T_1$	$T_2$
LOCK-X(A)	
READ(A)	
WRITE(A)	
UNLOCK(A)	
	LOCK-S(B)
	READ(B)
	UNLOCK(B)
LOCK-X(B)	
READ(B)	
WRITE(B)	
UNLOCK(B)	
	LOCK-S(A)
	READ(A)
	UNLOCK(A)

## Two phase locking protocol(2PL)

- The protocol ensures that each transaction issue lock and unlock requests in two phases, note that each transaction will be 2 phased not schedule.
- **Growing phase**- A transaction may obtain locks, but not release any locks.
- **Shrinking phase**- A transaction may release locks, but may not obtain any new locks.



$T_1$	$T_2$
LOCK-X(A)	
READ(A)	
WRITE(A)	
UNLOCK(A)	
	LOCK-S(B)
	READ(B)
	UNLOCK(B)
LOCK-X(B)	
READ(B)	
WRITE(B)	
UNLOCK(B)	
	LOCK-S(A)
	READ(A)
	UNLOCK(B)
	UNLOCK(A)
	UNLOCK(A)

$T_1$	$T_2$
LOCK-X(A)	
READ(A)	
WRITE(A)	
	LOCK-S(B)
	READ(B)
LOCK-X(B)	
READ(B)	
WRITE(B)	
	LOCK-S(A)
	READ(A)
	UNLOCK(B)
	UNLOCK(A)
	UNLOCK(B)
	UNLOCK(A)

Deadlock Problem

## Properties

- 2PL ensures conflict serializability, and the ordering of transactions over lock points is itself a serializability order of a schedule in 2PL.
- If a schedule is allowed in the 2PL protocol, then it is definitely always conflict serializable. But it is not necessary that if a schedule is conflict serializable, then it will be generated by 2PL. The equivalent serial schedule is based on the order of lock points.
- View serializability is also guaranteed.
- Does not ensure freedom from deadlock
- May cause non-recoverability.
- Cascading rollback may occur.

	Conflict Serializability	View Serializability	Recoverability	Cascadelessness	Deadlock Freedom
Time Stamp Ordering	YES	YES	NO	NO	YES
Thomas Write Rule	NO	YES	NO	NO	YES
Basic 2PL	YES	YES	NO	NO	NO
Conservative 2PL					
Rigorous 2PL					
Strict 2PL					

## Conservative 2PL

- The idea is that there is no growing phase; transactions start directly from the lock point, i.e., transactions must first acquire all the required locks, and then only can start execution. If all the locks are not available, then the transaction must release the acquired locks and must wait.
- The shrinking phase will work as usual, and transactions can unlock any data item anytime.
- We must have knowledge of the future to understand what data is required so that we can use it.

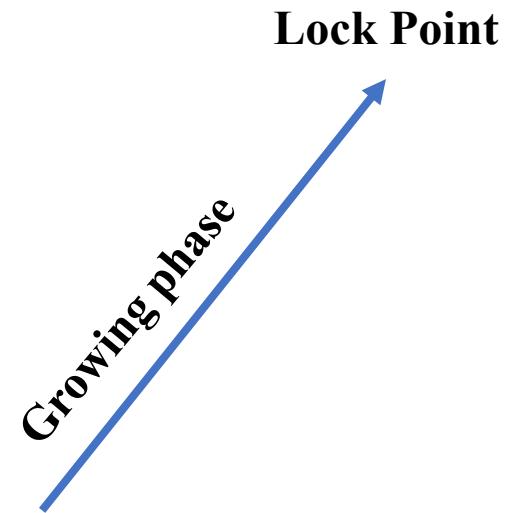
**Lock Point**

shrinking phase

Conservative 2PL

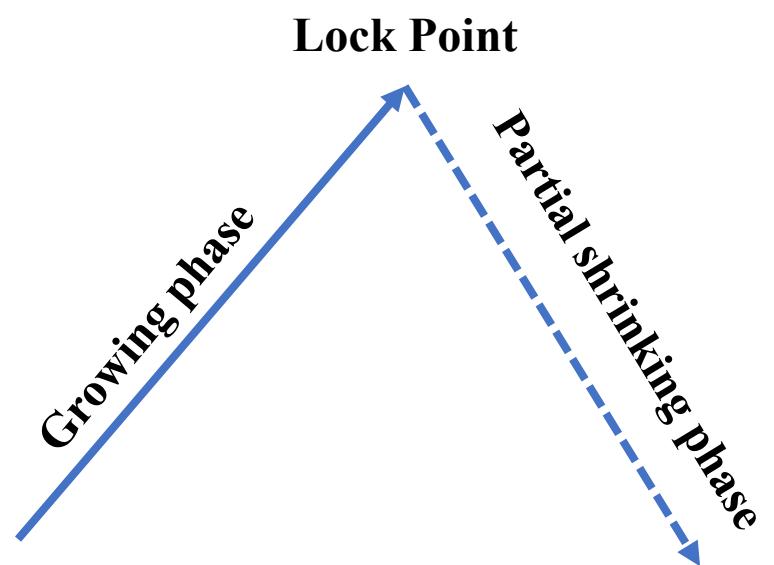
## RIGOROUS 2PL

- ❖ Requires that all locks be held until the transaction commits.
- ❖ This protocol requires that locking be two phase and also all the locks taken be held by transaction until that transaction commit.
- ❖ Hence, there is no shrinking phase in the system.



## STRICT 2PL

- ❑ That all exclusive-mode locks taken by a transaction be held until that transaction commits. This requirement ensures that any data written by an uncommitted transaction is locked in exclusive mode until the transaction commits, preventing any other transaction from reading the data.
- ❑ This protocol requires that locking be two-phase and also that exclusive-mode locks taken by a transaction be held until that transaction commits.
- ❑ So it is a simplified form of rigorous 2PL



	Conflict Serializability	View Serializability	Recoverability	Cascadelessness	Deadlock Freedom
Time Stamp Ordering	YES	YES	NO	NO	YES
Thomas Write Rule	NO	YES	NO	NO	YES
Basic 2PL	YES	YES	NO	NO	NO
Conservative 2PL	YES	YES	NO	NO	YES
Rigorous 2PL	YES	YES	YES	YES	NO
Strict 2PL	YES	YES	YES	YES	NO

## Validation-Based Protocols

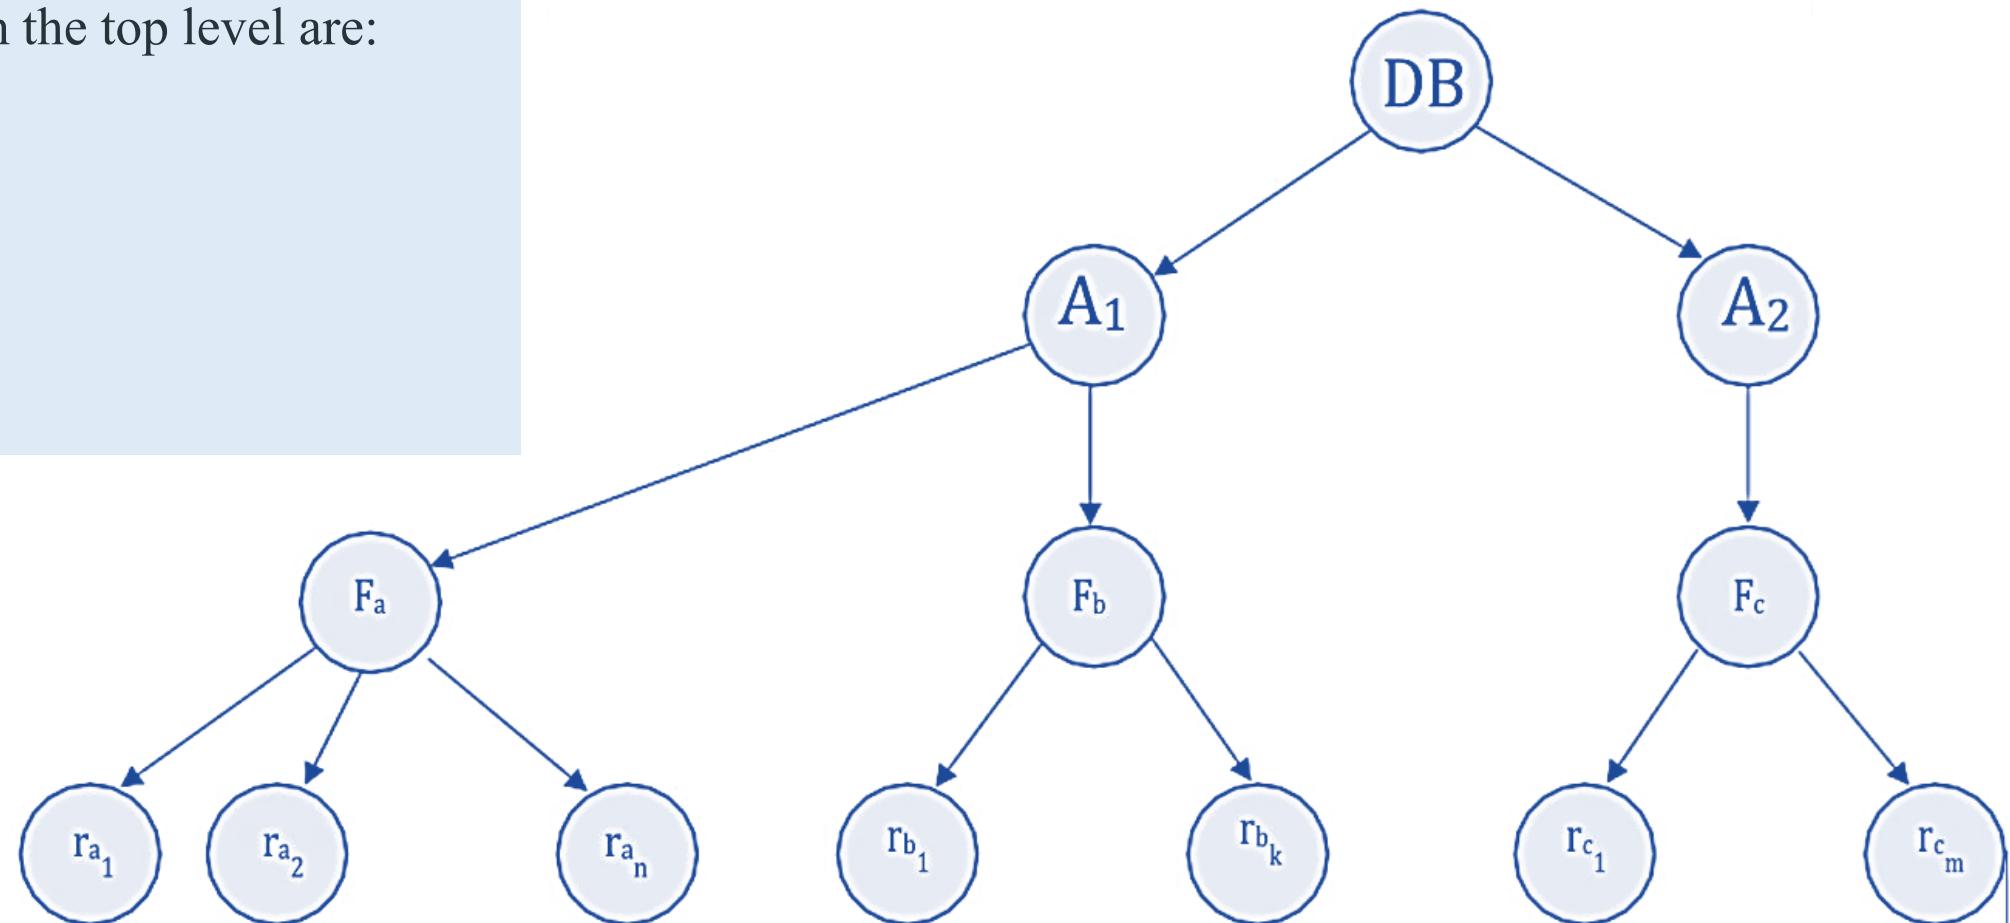
- ❑ In cases where a majority of transactions are read-only transactions, the rate of conflicts among transactions may be low.
- ❑ Thus, many of these transactions, if executed without the supervision of a concurrency-control scheme, would nevertheless leave the system in a consistent state.
- ❑ A concurrency-control scheme imposes overhead of code execution and possible delay of transactions. It may be better to use an alternative scheme that imposes less overhead.
- ❑ A difficulty in reducing the overhead is that we do not know in advance which transactions will be involved in a conflict. To gain that knowledge, we need a scheme for *monitoring* the system.
- ❑ The **validation protocol** requires that each transaction  $T_i$  executes in two or three different phases in its lifetime, depending on whether it is a read-only or an update transaction. The phases are, in order:
  - ❑ **Read phase.** During this phase, the system executes transaction  $T_i$ . It reads the values of the various data items and stores them in variables local to  $T_i$ . It performs all write operations on temporary local variables, without updates of the actual database.
  - ❑ **Validation phase.** The validation test (described below) is applied to transaction  $T_i$ . This determines whether  $T_i$  is allowed to proceed to the write phase without causing a violation of serializability. If a transaction fails the validation test, the system aborts the transaction.
  - ❑ **Write phase.** If the validation test succeeds for transaction  $T_i$ , the temporary local variables that hold the results of any write operations performed by  $T_i$  are copied to the database. Read-only transactions omit this phase.

## Multiple Granularity

It is the size of the data item allowed to lock. Now *Multiple Granularity* means hierarchically breaking up the database into blocks that can be locked and can be tracked what needs to lock and in what fashion. Such a hierarchy can be represented graphically as a tree.

The levels starting from the top level are:

- database
- area
- file
- record



In the 2-phase locking protocol, it shall use shared and exclusive lock modes. When a transaction locks a node, in either shared or exclusive mode, the transaction also implicitly locks all the descendants of that node in the same lock mode.

## Intention Mode Lock

In addition to **S** and **X** lock modes, there are three additional lock modes with multiple granularities:

- Intention-Shared (IS):** explicit locking at a lower level of the tree but only with shared locks.
- Intention-Exclusive (IX):** explicit locking at a lower level with exclusive or shared locks.
- Shared & Intention-Exclusive (SIX):** the subtree rooted by that node is locked explicitly in shared mode and explicit locking is being done at a lower level with exclusive mode locks.

	IS	IX	S	SIX	X
IS	✓	✓	✓	✓	✗
IX	✓	✓	✗	✗	✗
S	✓	✗	✓	✗	✗
SIX	✓	✗	✗	✗	✗
X	✗	✗	✗	✗	✗

IS : Intention Shared

IX : Intention Exclusive

S : Shared

X : Exclusive

SIX : Shared & Intention Exclusive

**compatibility matrix for these lock modes**

The multiple-granularity locking protocol uses the intention lock modes to ensure serializability. Multiple-granularity protocol requires that locks be acquired in top-down (root-to-leaf) order, whereas locks must be released in bottom-up (leaf-to-root) order.

- Database objects can be locked by transactions at various granularities using MGL, ranging from individual data points to whole tables or databases.
- Because multiple granularity locking shortens the time that transactions must wait for locks to be released, it can increase concurrency.
- When compared to other concurrency control algorithms, such two-phase locking (2PL), MGL can be more difficult to implement and maintain.
- This protocol enhances concurrency and reduces lock overhead. Deadlock is still possible in the multiple-granularity protocol, as it is in the two-phase locking protocol.

# Types of Database Failures

There are many types of failures that can affect database processing. Some failures affect the main memory only, while others involve secondary storage. Following are the types of failure:

- **Hardware failures:** Hardware failures may include memory errors, disk crashes, bad disk sectors, disk full errors and so on. Hardware failures can also be attributed to design errors, inadequate (poor) quality control during fabrication, overloading (use of under-capacity components) and wearout of mechanical parts.
- **Software failures:** Software failures may include failures related to software such as, operating system, DBMS software, application programs and so on.
- **System crashes:** System crashes are due to hardware or software .

# Database Recovery

is a process of recovering or restoring data in the database when a data loss occurs or data gets deleted by system crash, hacking, errors in the transaction, damage occurred coincidentally, by viruses, sudden terrible failure, commands incorrect implementation, etc. Data loss or failures happen in databases like other systems but the data stored in the database should be available whenever it's required. For fast restoration or recovery of data, the database must hold tools which recover the data efficiently. It should have atomicity means either the transactions showing the consequence of successful accomplishment perpetually in the database or the transaction must have no sign of accomplishment consequence in the database.

From any failure set of circumstances, there are both voluntary and involuntary ways for both, backing up of data and recovery. So, recovery techniques which are based on deferred update and immediate update or backing up data can be used to stop loss in the database.

## Crash recovery

**Crash recovery** is the operation through which the database is transferred back to a compatible and operational condition. In DBMS, this is performed by rolling back insufficient transactions and finishing perpetrated transactions that even now existed in memory when the crash took place.

With many transactions being implemented with each second shows that, DBMS may be a tremendously complex system. The fundamental hardware of the system manages to sustain robustness and stiffness of software which depends upon its complex design. It's anticipated that the system would go behind with some methodology or techniques to restore lost data when it fails or crashes in between the transactions.

# Classification of failure

The following points are the generalization of failure into various classifications, to examine the source of a problem,

1. **Transaction failure:** a transaction must terminate when it arrives at a point from where it can't extend any further and when it fails to implement the operation.  
Transaction failure reasons could be,
  - **Logical errors:** The errors which take place in some code or any intrinsic error situation, where a transaction cannot properly fulfil.
  - **System errors:** The errors which take place when the database management system is not able to implement the active transaction, or it has to terminate it because of some conditions in a system.
2. **System Crash:** There are issues which may stop the system unexpectedly from outside and may create the system condition to crash. For example, disturbance or interference in the power supply may create the system condition of fundamental hardware or software to crash or failure.
3. **Disk Failure:** Disk failures comprise bad sectors evolution in the disk, disk inaccessibility, and head crash in the disk, other failures which damage disk storage completely or its particular parts.

# Recovery and Atomicity

To recover and also to sustain the transaction atomicity, there are two types of methodology,

- Sustaining each transaction logs and before actually improving the database put them down onto some storage which is substantial.
- Sustaining shadow paging, in which on a volatile memory the improvements are completed and afterward, the real database is reformed.

## Log-based Recovery

The log is an order of sequence of records, which sustains the operations record accomplished by a transaction in the database. Before the specific changes and improvements survive on a storage media which is stable and failing securely, it's essential that the logs area unit put down in storage.

Following are the workings of Log-based Recovery,

The log file is not damaged on a stable storage media.

Log-based recovery puts down a log regarding a transaction when a transaction begins to be involved in the system and starts implementation.

# **Recovery with Concurrent Transactions**

The logs are interleaved when multiple transactions are being implemented in collateral. It would be difficult for the system of recovery to make an order of sequence of all logs again, and then start recovering at the time of recovery. Most recent times Database systems use the abstraction of 'checkpoints' to make this condition uncomplicated.

## **Checkpoint**

The checkpoint is an established process where all the logs which are previously used are clear out from the system and stored perpetually in a storage disk. Checkpoint mention a point before which the DBMS was in a compatible state, and all the transactions were perpetrated.